

ZeBu[®] Simulation Acceleration User Guide

Version V-2024.03-1, July 2024



Copyright and Proprietary Information Notice

© 2024 Synopsys, Inc. This Synopsys software and all associated documentation are proprietary to Synopsys, Inc. and may only be used pursuant to the terms and conditions of a written license agreement with Synopsys, Inc. All other use, reproduction, modification, or distribution of the Synopsys software or the associated documentation is strictly prohibited.

Destination Control Statement

All technical data contained in this publication is subject to the export control laws of the United States of America. Disclosure to nationals of other countries contrary to United States law is prohibited. It is the reader's responsibility to determine the applicable regulations and to comply with them.

Disclaimer

SYNOPSYS, INC., AND ITS LICENSORS MAKE NO WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, WITH REGARD TO THIS MATERIAL, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE.

Trademarks

Synopsys and certain Synopsys product names are trademarks of Synopsys, as set forth at <https://www.synopsys.com/company/legal/trademarks-brands.html>. All other product or company names may be trademarks of their respective owners.

Free and Open-Source Licensing Notices

If applicable, Free and Open-Source Software (FOSS) licensing notices are available in the product installation.

Third-Party Links

Any links to third-party websites included in this document are for your convenience only. Synopsys does not endorse and is not responsible for such websites and their practices, including privacy practices, availability, and content.

www.synopsys.com

Contents

About This Book	9
Intended Audience	9
Contents of This Book	9
Related Documentation	10
Typographical Conventions	12
Synopsys Statement on Inclusivity and Diversity	12
<hr/>	
1. Introduction to Simulation Acceleration Cosimulation	14
Simulation Acceleration Flow	15
V2VX Mode	15
V2Z Mode	15
Simulation Acceleration Compile and Runtime Options	17
V2VX Mode Compilation and Runtime Options	17
V2Z Mode Compilation and Runtime Options	17
VCS Compile-Time Option to Generate Additional Diagnostic Data	19
On-Screen Messages	20
Simulation Acceleration Coding Guidelines	21
SystemC Support in Simulation Acceleration	22
SystemC Limitations	23
Other Supported Features in Simulation Acceleration	23
Interface Appearing in Both SW and HW	24
Moving a HW Module/Interface/Task/Function/Block to Execute in SW	24
Using Compile-Time Options	24
Using an Embedded Source Code SystemVerilog Attribute, Comment Pragma, or UTF Command	25
Moving Initial Blocks From HW to SW Automatically	26
SVA Assertions are Supported in SW and HW That Is Not Mapped to IF FPGAs	29
VHDL Support and Limitations	29
VHDL Support	29
Limitations	31
Moving Glue Logic to Improve Performance	32
Types of Glue logic Supported	32

Contents

Moving Glue Logic	32
Support for Final Block in Hardware	33
Limitation	34
Support for Concurrent Assertion in Hardware	34
Enabling Concurrent Assertion in Hardware	34
Example	34
Limitations	35
Moving dumpvars/dumpports to the DUT Automatically	35
Example 1: Dedicated Initial Block	36
Example 2: Non-Dedicated Initial Block	36
Limitations	37
Miscellaneous Features	37
Other Limitations That Do Not Include Signals or Subroutines	38
Temporary Known Simulation Acceleration Limitations	41
<hr/>	
2. Supported Syntax in Simulation Acceleration	42
Supported Syntax for Using Signals as the Communication Mechanism	42
Signal Communication Using Primitive Data Types	43
Signal Communication Using Aggregate Data Types	43
Example	44
Signal Communication Using Interface Ports	44
Example	44
Signal Communication Using Real Types	45
Use Model	45
Limitations	45
Signal Communication Using XMRs	45
Optimizing Performance for Clock XMRs From SW to HW (clkopt)	47
Support for iff Construct with clkopt	47
Enhancements to clkopt Optimization	48
Optimizing Performance for XMRs With Force and Release Statements	48
SW Accessing HW Memory Using \$readmem and \$writemem System	
Tasks	48
Example	49
Accessing and Forcing HW Signals From Within SW Using \$hdl_xmr and	
\$hdl_xmr_force	49
\$hdl_xmr: Access HW Signals From Within SW	49
\$hdl_xmr_force: Force HW Signals From Within SW	50
Signal Communication Through External Subroutines	50
Example 1	51
Example 2	51

Contents

Strength Analysis in Simulation Acceleration	51
Enabling Strength Analysis	52
Viewing Strength Results in HW	52
Limitations With Signal-Based Communication	52
Supported Syntax for Using Subroutines as the Communication Mechanism	54
Interface Task/Function (Subroutines) Calls From SW to HW	55
Class Method Calls From HW Interfaces to SW	56
Example	56
DPI Subroutine Call From HW Interfaces	57
Streaming Export Methods	57
UVM Messaging Support	57
Import Calling an Export	59
Example	59
\$test\$plusargs, \$value\$plusargs Support	60
Limitation	60
System Tasks	60
System Task Support for \$display	61
Enabling \$display Optimization for All HW Modules	62
Enabling \$display Optimization for Specific HW Modules	62
Enabling \$display Optimization for Backdoor Communication	62
Pragma Support	63
UVM Report Support	63
Limitations of Display Optimization	64
Enabling Concurrent Export Task/Function Calls	64
Use Model	64
Supported and Unsupported Behavioral Compiler Syntax	65
Supported Behavioral Compiler Syntax in Simulation Acceleration	65
#delay Support	66
Clocking Edge in for Loops	67
Multiple Clocks or Edge Expressions in the Same Process	67
Complex Edge Combination	67
Combinations of Edge and Level Events	67
Bounded Loops	67
Unbounded Loops	68
Wait Statements	68
Named Events	68
Behavioral Compiler Supports all SystemVerilog Data Types	68
Clocking Blocks	68
Fork/Join	69
SystemVerilog Strings, \$sformatf, %m are Supported	69
Usage of Memory in Behavioral Code	70

Unsupported Behavioral Compiler Syntax in Simulation Acceleration	72
<hr/>	
3. Scheduling Semantics and Race Conditions	74
Initialization Between SW and HW	75
Signal Handshake Between SW and HW	75
Resolving Clock/Data Race Conditions Between SW and HW	75
Subroutine Handshake Between SW and HW	76
Import Calls and Non-Blocking/Blocking Assignments	76
Export Calls Scheduling	77
<hr/>	
4. Incremental Compile	78
Modes for Performing Incremental Compile with zCui	79
Support for \$hw_force ECO	80
Prerequisites	80
Use Model	80
Limitations	81
Key Messages in Simulation Acceleration Incremental Compile	81
No Warnings Reported When Generating simv	81
Warnings Reported While Generating simv	82
Communication Changes Caused Warnings While Generating simv	82
New simv Cannot be Generated	82
<hr/>	
5. SVA Assertion and Coverage Support	83
Enabling Functional Coverage and Assertions at Runtime	83
Controlling SVAs at Runtime	84
<hr/>	
6. Generating the Profiler Report	86
Enabling Profiler	86
Profiler Report Sections	87
Example: Communication Profiler Report	88
<hr/>	
7. Debugging and Reports	95
Debug Options for Communication Mechanisms	95

Contents

Functional Debug in Simulation Acceleration	96
Enabling UCLI Signal Access for ZeBu Compile	97
Enabling UCLI Waveform Output for ZeBu Compile	97
Supported UCLI commands in Simulation Acceleration	98
get UCLI Command	98
memory UCLI Command	99
stop UCLI Command	100
run UCLI Command	102
force, release, show UCLI Commands	102
restart UCLI Command	103
System Information Commands	104
dump Commands	104
sniffer Command	105
powermgt UCLI Command	107
Verdi Interactive Debug for Simulation Acceleration	108
Waveform Outputting Using UCLI	108
Dynamic-Probe	109
FWC Probe	110
QiWC Probe	110
Outputting Signals of HW Module Instances That Were Moved to SW	111
Multi-ZTDB Output, Multi-ValueSet Capture	111
Output FWC/QiWC Waveform Only on Time Advance	112
Using Verdi for Debug	112
Output ZeBu XTOR Internal Signals	112
Waveform Expansion in Simulation Acceleration	113
Waveform Outputting config slice for zSimzilla	114
Using simprofile with Simulation Acceleration	115
Viewing Timeline Using simprofile	115
Viewing Delay Events Using simprofile	116
<hr/>	
8. Timing Analysis	118
UTF File Changes	118
Viewing the zTime Report for Critical Output Routing Data Paths	119
Analyzing a Critical Path Using zTime	121
<hr/>	
9. Using Time Decoupled Mode for Performance Improvements	123
Consideration for Using Time Decoupling	123
Use Model for the Time Decouple Mode	124
Compile-Time Use Model	124

Contents

Automatically Infer All Input-Only Functions	125
Manually Specify Input-Only Functions	125
Runtime Use Model	125
Limitations of Time Decouple Mode	126

10. Enabling zDPI in Simulation Acceleration	127
Use Model for Enabling zDPI in the V2Z Flow	127
Compile-Time Use Model	127
Runtime Use Model	128
Use Model for Enabling zDPI in the V2VX Flow	128
Support to Call Export DPI Function Calls	128
Limitations of zDPI Support in Simulation Acceleration	129

11. Estimating DUT Power Requirements Using Power Profiling	130
Enabling Weighted Activity Profile (WAP) Power Profiling	130
Use Model of the wap Command	130
Sample UCLI Script Using the WAP Command Specifications	131

Preface

This chapter has the following sections:

- [About This Book](#)
- [Intended Audience](#)
- [Contents of This Book](#)
- [Related Documentation](#)
- [Typographical Conventions](#)
- [Synopsys Statement on Inclusivity and Diversity](#)

About This Book

The *Simulation Acceleration User Guide* describes how to use Simulation Acceleration to enable cosimulating SystemVerilog testbenches with the DUT without any changes to the code. Simulation Acceleration enables communicating between software (SW) and hardware (HW) parts using SystemVerilog mechanisms, such as ports on a module or interface and calling of subroutines implemented in an interface.

Intended Audience

This guide is written for design engineers to assist them with using Simulation Acceleration in ZeBu.

These engineers should have knowledge of the following Synopsys tools:

- ZeBu Server
- VCS
- Verdi

Contents of This Book

This document has the following sections.

Section	Describes
Introduction to Simulation Acceleration Cosimulation	Steps to deploy Simulation Acceleration cosimulation flow (V2VX and V2Z modes), Simulation Acceleration Compile and Runtime Options, Simulation Acceleration Coding Guidelines, and SystemC Support in Simulation Acceleration
Supported Syntax in Simulation Acceleration	Supported syntax for using signals as the communication mechanism, supported syntax for using subroutines as the communication mechanism, and supported and unsupported behavioral compiler syntax
Scheduling Semantics and Race Conditions	Describes various topics, such as resolving Clock/Data race conditions between SW and HW
Incremental Compile	How to enable incremental compile on a given build
SVA Assertion and Coverage Support	How to enable functional coverage and assertions at compile time; Also, describes controlling SVAs at runtime
Generating the Profiler Report	The SimXL Profiler report generated by the Simulation Acceleration [Communication] Profiler tool for Simulation Acceleration.
Debugging and Reports	Debug options for communication mechanisms, functional debug in Simulation Acceleration, and using simprofile with Simulation Acceleration
Timing Analysis	Viewing the zTime report for critical output routing data paths and analyzing a critical path using zTime
Using Time Decoupled Mode for Performance Improvements	The use model to mark functions, which only have inputs, so that these functions are not monitored or controlled by Simulation Acceleration transactors. This is known as time decoupling of functions, which improves performance.
Enabling zDPI in Simulation Acceleration	Use model for enabling zDPI in V2Z and V2VX flows
Estimating DUT Power Requirements Using Power Profiling	The Power Profiling feature in Simulation Acceleration enables you to estimate the power requirements of the DUT, which resides in the hardware; this section describes how to enable WAP profiling

Related Documentation

Document Name	Description
<i>ZeBu User Guide</i>	Provides detailed information on using ZeBu.

Document Name	Description
<i>ZeBu Debug Guide</i>	Provides information on tools you can use for debugging.
<i>ZeBu Debug Methodology Guide</i>	Provides debug methodologies that you can use for debugging.
<i>ZeBu Unified Command-Line User Guide</i>	Provides the usage of Unified Command-Line Interface (UCLI) for debugging your design.
<i>ZeBu UTF Reference Guide</i>	Describes Unified Tcl Format (UTF) commands used with ZeBu.
<i>ZeBu Power Aware Verification User Guide</i>	Describes how to use Power Aware verification in ZeBu environment, from the source files to runtime.
<i>ZeBu Functional Coverage User Guide</i>	Describes collecting functional coverage in emulation.
<i>Simulation Acceleration User Guide</i>	Provides information on how to use Simulation Acceleration to enable cosimulating SystemVerilog testbenches with the DUT
<i>ZeBu Verdi Integration Guide</i>	Provides Verdi features that you can use with ZeBu. This document is available in the Verdi documentation set.
<i>ZeBu Runtime Performance Analysis With zTune User Guide</i>	Provides information about runtime emulation performance analysis with zTune.
<i>ZeBu Custom DPI Based Transactors User Guide</i>	Describes ZEMI-3 that enables writing transactors for functional testing of a design.
<i>ZeBu LCA Features Guide</i>	Provides a list of Limited Customer Availability (LCA) features available with ZeBu.
<i>ZeBu Synthesis Verification User Guide</i>	Provides a description of zFmCheck.
<i>ZeBu Transactors Compilation Application Note</i>	Provides detailed steps to instantiate and compile a ZeBu transactor.
<i>ZeBu zManualPartitioner Application Note</i>	Describes the <code>zManualPartitioner</code> feature for ZeBu. It is a graphical interface to manually partition a design.
<i>ZeBu Hybrid Emulation Application Note</i>	Provides an overview of the hybrid emulation solution and its components.

Typographical Conventions

This document uses the following typographical conventions:

To indicate	Convention Used
Program code	<code>OUT <= IN;</code>
Object names	<code>OUT</code>
Variables representing objects names	<code><sig-name></code>
Message	Active low signal name ' <code><sig-name></code> ' must end with <code>_X</code> .
Message location	<code>OUT <= IN;</code>
Example with message removed	<code>OUT_X <= IN;</code>
Important Information	NOTE: This rule...

The following table describes the syntax used in this document:

Syntax	Description
<code>[]</code> (Square brackets)	An optional entry
<code>{ }</code> (Curly braces)	An entry that can be specified one time or multiple times
<code> </code> (Vertical bar)	A list of choices out of which you can choose one
<code>...</code> (Horizontal ellipsis)	Other options that you can specify

Synopsys Statement on Inclusivity and Diversity

Synopsys is committed to creating an inclusive environment where every employee, customer, and partner feels welcomed. We are reviewing and removing exclusionary language from our products and supporting customer-facing collateral. Our effort also includes internal initiatives to remove biased language from our engineering and working environment, including terms that are embedded in our software and IPs. At the same time, we are working to ensure that our web content and software applications are usable to people of varying abilities. You may still find examples of non-inclusive language in our

software or documentation as our IPs implement industry-standard specifications that are currently under review to remove exclusionary language.

1

Introduction to Simulation Acceleration Cosimulation

Simulation Acceleration enables cosimulating SystemVerilog testbenches with the DUT without any changes to the code.

The usage of Simulation Acceleration in ZeBu has the following advantages:

- It provides a mechanism for easy creation of a generic simulation testbench setup that can be easily ported to the ZeBu emulation platform.
- It provides a well-structured approach as recommended by standard and advanced SystemVerilog based verification methodologies, such as UVM.
- It provides an incremental approach to bringing up a working simulation environment into an emulation environment. The approach starts with a signal-based communication mechanism between SW (VCS) and HW (ZeBu). Then, progresses with working through the different testbench interfaces toward a transaction-based communication mechanism.

Simulation Acceleration enables communicating between software (SW) and hardware (HW) parts using SystemVerilog mechanisms, such as ports on a module or interface and calling of subroutines implemented in an interface.

For more information on Simulation Acceleration, see the following topics:

- [Simulation Acceleration Flow](#)
- [Simulation Acceleration Compile and Runtime Options](#)
- [Simulation Acceleration Coding Guidelines](#)
- [SystemC Support in Simulation Acceleration](#)
- [Other Supported Features in Simulation Acceleration](#)

Simulation Acceleration Flow

You can start to deploy the Simulation Acceleration cosimulation flow on a working VCS test environment. The following modes are used:

- [V2VX Mode](#)
- [V2Z Mode](#)

V2VX Mode

The V2VX mode is used to compile and run a test case in the environment using VCS only to validate the readiness of environment for Simulation Acceleration.

- If the compile fails, there could be unsupported syntax used. In this case the compiler provides a compilation Not Yet Implemented (NYI) error message, indicating the specific syntax that is not supported.
- If running the test fails, there could be multiple reasons for failure, such as the following:
 - Race conditions between HW and SW that were not exposed in VCS.
 - Incorrect Simulation Acceleration Transformations and Mapping.
 - Usage of unsupported features.

If the test fails, it is recommended to do a waveform comparison between a regular VCS run and a V2VX VCS run, using Verdi nCompare capabilities.

V2Z Mode

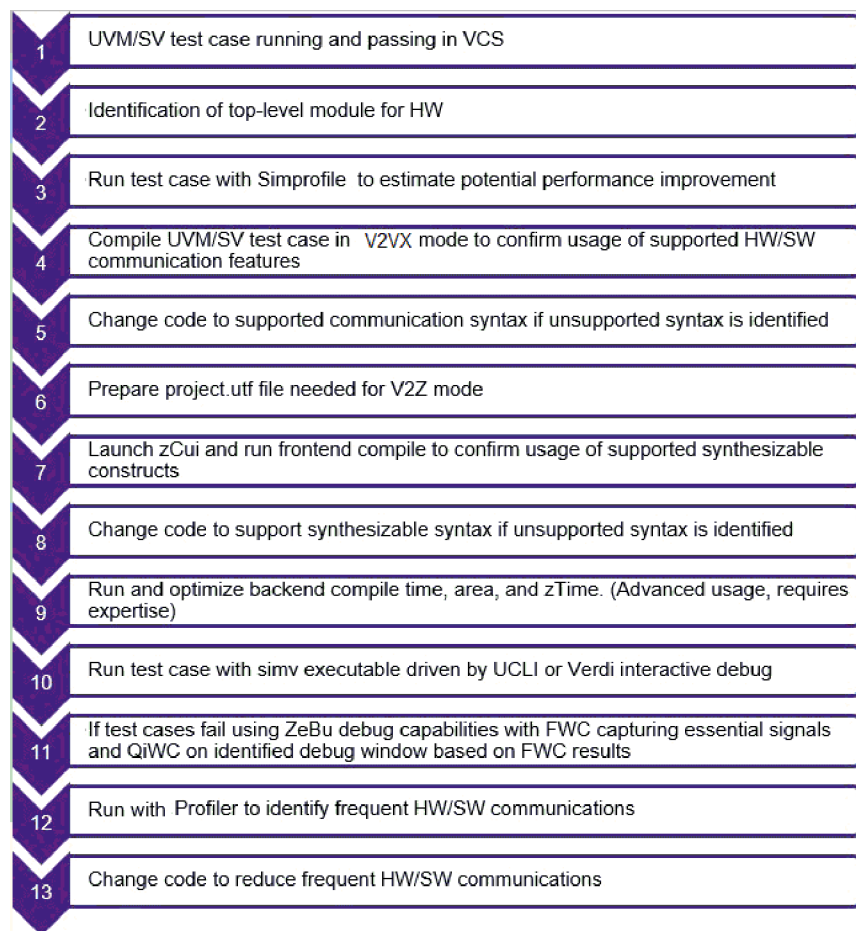
The V2Z mode is used to compile and run a test case in a cosimulation environment. SW is compiled with VCS and HW is compiled using ZeBu.

- If the compile fails, there could be multiple reasons for failure, such as the following:
 - Incomplete port connectivity to HW
 - Synthesis issues caused by usage of non-synthesizable syntax
 - Place-and-Route issues
- If running the test fails, there could be multiple reasons for failure, such as the following:
 - Race conditions between HW and SW that were not exposed in VCS
 - Incorrect Simulation Acceleration Transformations and Mapping

- Incorrect RTL code generated by Behavioral Compiler.
- Usage of unsupported features.
- Non congruence between simulation and emulation, for example emulator does not support 'z' and 'x' values.
- If the test fails, it is recommended to do a waveform comparison between a regular VCS run and a V2Z run, using Verdi nCompare capabilities

The `simv` executable is used to run the Simulation Acceleration environment to generate waveforms and log files.

Figure 1 *Simulation Acceleration Flow*



Simulation Acceleration Compile and Runtime Options

The options for compilation and runtime depend on the mode. For more information, see the following subsections:

- [V2VX Mode Compilation and Runtime Options](#)
- [V2Z Mode Compilation and Runtime Options](#)
- [VCS Compile-Time Option to Generate Additional Diagnostic Data](#)
- [On-Screen Messages](#)

Note:

It is recommended not to run different compilations parallelly in the same directory. The content in the `csrc` directory might be corrupted and the VCS compilation might behave in an unexpected manner.

V2VX Mode Compilation and Runtime Options

- Uses only VCS for both HW and SW compilation. Used for pipe cleaning before moving to ZeBu compile.
- To enable the V2VX mode, use the following VCS Compile option:
`-Xhwcossimtest=v2vx`

- To specify HW top-level instance, use the following VCS Compile option:

`-Xhdl_cosim_dut <hierarchical path to the DUT instance>`

If the module is compiled as a top-level module, it is the name of the module.

- To specify that all instances of a module are HW top-level instance, use the following VCS Compile option: `-Xhdl_cosim_dut_module <module-name>`
- Run using existing `simv` options.

V2Z Mode Compilation and Runtime Options

- Uses ZeBu for HW compilation and VCS for SW compilation.
- The ZeBu **zCui** tool does the full compile with a UTF file as input and the UTF file pointing to the VCS compile command using the `vcs_exec_command` UTF command.

- For compile, add the following UTF options and VCS compile options:
 - UTF command to enable Simulation Acceleration: `simxl -enable true`
 - UTF command to specify HW top-level instance when an instance path is required to select a specific instance subtree: `simxl_set_hwtop -instance <hierarchical path to the DUT instance>`. All logic from that instance and below goes into HW. If the module is compiled as a top-level module, the `-module` option below is required.
 - The UTF command to specify that all instances of a module are HW top-level instances is as follows: `simxl_set_hwtop -module <module-name>`

Use this command if the HW top is a single instance module.

Multiple `simxl_set_hwtop` (respectively for `v2vx`, `-Xhdl_cosim_dut`, `-Xhdl_cosim_dut_module`) commands can be used to specify more than one top, the following restrictions are applicable to both V2VX and V2Z modes:

- All HW top levels need to be instantiated in the same module.
- For example, specifying `"top1.dut1"` and `"top2.dut2"` as dut-instances would give an NYI error because `dut1` and `dut2` are in two different top modules, `top1` and `top2`.
- A HW top module cannot be instantiated in another HW top module.

For example, suppose `Top` has instances of module `"A"` and `"B"`, named `instA` and `instB`. Module `"A"` has again an instance of module `"B"`. In this case, you cannot specify both `"top.instA"` and `"top.instB"` as HW top instances.

- A HW top instance cannot be within a hierarchy that has been moved to SW.
- A HW top instance cannot be an instance-array of size more than 1. For example:

```
HWTOP dut_inst[1:0]();
```

If you specify `"HWTOP"` as `hwtop` module, an NYI error is reported.

Usage of `$dumpvars` for specifying waveform dumping directives can be used in any of the HW top modules. It is also possible to have a `$dumpvars` HW top module (For example, `dumpvars_module.sv`) to specify waveform dumping directives for any of the signals/instances in the instance subtrees of the other HW top modules.

- To execute the ZeBu compile, perform the following:
 - Open the **zCui** compilation GUI using the following:

```
zCui -u project.utf
```
 - Click the **Make Target** button in the **zCui** GUI to start the compile process.

- In the case where compilation of the **zCui** front end needs to be exported to enable **zCui** back-end compile at a different location, perform the following:
 - Prior to the front-end compile, run the following command to set the `ZCUI_EXPERT` environment variable:

```
setenv ZCUI_EXPERT exportFrontEnd
```
 - Execute the front-end compile only using **zCui**.
 - After front-end compile, move the tar or zip file, `frontend.tgz`, to the different location and extract.
 - Run the back-end compile using the following command and options:

```
zCui --rf feBeSplit/frontendRestore.xcui
```

VCS Compile-Time Option to Generate Additional Diagnostic Data

VCS additional compile-time option to generate diagnostic data in the current run directory:

```
-simxl=diag
```

To generated VCS compile diagnostic file at the VCS stage, which is the early stage of **zCui** compilation, specify the following:

```
-simxl=comp_diag
```

The `simxlCompDiag.txt` file is generated. This file contains the following:

- All Top level modules including TB
- All DUT top modules
- All embedded TB modules (ETB)
- VCS command line: Switches used for VCS elaboration
- Simulation Acceleration interface signals
 - List of INPUT signal full instance name exported with bit width and module/interface name. In the end, total input signals exported with total bit width.
 - List of OUTPUT signal full instance name exported with bit width and module/interface name. In the end total output signals exported with total bit width.
- Simulation Acceleration task functions

- List of exported tasks (full task scope) with detailed argument list (bit width, direction, name) and its module name.
- List of imported tasks (HW module caller instance) with detailed argument list (bit width, direction, name) and its module name.
- Simulation Acceleration zcei message port cost: Rough estimate of "In/Out" port message port cost per instance
- Transactor name and full instance path: All the transactor name/hierarchy

This report does not contain:

- Transactor information is dumped at VCS side only in NO_DVE_FLOW. For other flow, transactor information is dumped at ZEBU side. Since we only track VCS side diagnostics, the transactor information will not show up consistently in `simx1CompDiag.txt` file.
- Only embedded TB modules (ETB) are reported. TF/blocks are not covered.

On-Screen Messages

When you run `simv`, by default, only critical information, such as the following, is displayed on-screen (in terminal):

- Testbench messages
- Necessary ZeBu messages, such as log file location and license checks status
- Error messages and fatal messages

To view all information, see the log files.

If you need all information to be displayed on-screen, set the `ZEBU_DISABLE_PRINT` environment variable to any value other than `OK`, such as the following:

```
setenv ZEBU_DISABLE_PRINT 0
```

By default, the `ZEBU_DISABLE_PRINT` environment variable is to `OK` and therefore, only critical messages are displayed on-screen.

Simulation Acceleration Coding Guidelines

The following is a set of guidelines used to reduce the number of issues related to compile, runtime, and performance:

- Use supported coding styles mentioned in this document.
- Do not use unsupported coding styles mentioned in this document.
- Never instantiate a SW module in the HW side, use SystemVerilog bind for associating a SW module to HW in a modular fashion. You can then use compile-time options, attributes, or UTF commands to move the SW module to execute in SW.
- Do not use `#-delays` in code residing in the SW side, use clocks coming from the HW side to perform synchronization.
- Make sure to use single timescale across HW and SW.
- Avoid using precision units in the timescale as much as possible.
- Make sure all clocks driving HW are implemented in HW using RTL clocks only.
- When implementing RTL clocks using `#-delay` and procedural code, make sure that the procedural code does not contain mixed in implementations for other features. The tool requires RTL clock implementations to be standalone within its own procedural block.
- Try to use clocks derived from a single fastest clock as much as possible. There are techniques to do this even for fine-grained clock frequencies. Consult with Synopsys AEs.
- It is possible to define utility tasks, functions, and DPI in the HW top-level wrapper. The tool supports calling import methods from the HW top-level wrapper, which utility tasks, functions, and DPI may need.
- Avoid doing excessive backdoor writes to non-memory sequential elements because the mechanism that does backdoor writes is extremely slow in ZeBu. It is preferred to do front door writes.
- Do not initialize memory arrays using procedural SV code because memory arrays are initialized to zero by default when running on ZeBu. Use `+vcs+initreg+random` when compiling with VCS and `+vcs+initreg+0` when running with `simv` to achieve similar behavior for simulation only.
- Always baseline functionality with current step before moving to next steps. See steps mentioned in [Introduction to Simulation Acceleration Cosimulation](#).
- Always bring-up a Signal-Based Acceleration (SBA) environment (with steps mentioned in [Introduction to Simulation Acceleration Cosimulation](#)) first before bring-

up of Transaction-Based Acceleration (TBA) environment. Signal-Based Acceleration (SBA), where the HW top is the actual top.

- Manage all synchronization from the HW side; avoid scheduling on the SW side to reduce the SW overhead.
- Avoid using behavioral code as much as possible. RTL code is preferable as it is easier to debug and has less transformations.
- Use zero-time consuming functions and tasks when communication between HW and SW as much as possible. Implement state machines to manage the transaction.
- Do not access elements in HW from your UVM tests, sequences, and higher-level components. Leave them as abstract as possible to avoid HW compiles when they change. Essentially HW elements should be referenced only from UVM drivers and monitors only.
- Do not use clocking blocks in your interfaces this does not work well for acceleration.
- Implement tasks and state machines that UVM driver uses in your interface rather than driving signals directly from your UVM driver.
- Implement monitor loops in HW interfaces that call import tasks and functions implemented in SW in the UVM Monitor.
- Avoid using SV code that requires signal strength resolution to resolve multidriver logic on the boundary between SW and HW. This includes SV constructs, such as `tran*`, and so on.
- Avoid using bidirectional ports on the boundary between HW and SW as much as possible.
- Avoid using a mixture of blocking and non-blocking statements in the same always block in the SimXL transactor.
- Avoid using same variable driven from multiple always blocks in the SimXL transactor.

SystemC Support in Simulation Acceleration

Pass the `-sysc=show_sc_main` option to the elaboration command. You are required to pass the DUT path that is not prefixed with `"sc_main"` because of the current use model of SystemC in the Simulation Acceleration flow.

SystemC Limitations

SystemC specific limitations are as follows:

- Inout ports of two HW top modules connected in SystemC might not work because the port connectivity might not get traced through SystemC.

Example

```
module dut1(inout p1);  
  assign p1 = en ? drv1: 'bz;  
endmodule  
  
module dut2(inout p2);  
  assign p2 = en ? drv2: 'bz;  
endmodule
```

- If ports dut1.p1 and dut2.p2 are connected to same highconn in SystemC, it is a multiple driver scenario. This can work only if the information is passed to ZeBu at compile-time (that these two are shorted). However, since connectivity cannot be traced in SystemC, this scenario would not work.

Note:

There is no warning/error message for the above scenario.

- DUT ports connected through SystemC would not be optimized (only impacts Performance, no functional issue).

Example

If output port of dut1 is connected to input port of another dut2, the value from the output port would go from HW to SystemC and then come back to SW through the input port. It only impacts performance, but not the functionality.

- Additional UTF commands to give directives on how HW subroutines are configured can be provided by identifying the interface that the HW subroutine resides in. For example, you can specify the following to reduce the size of the communication mechanism of subroutines in *axi_interface* to enable ZeBu compile to converge:

```
zemi3 -module {axi_interface} -max_out_port_width 512  
      -max_in_port_width 512
```

Other Supported Features in Simulation Acceleration

The other supported elements required for the Simulation Acceleration flow are as follows:

- [Interface Appearing in Both SW and HW](#)
- [Moving a HW Module/Interface/Task/Function/Block to Execute in SW](#)

- [SVA Assertions are Supported in SW and HW That Is Not Mapped to IF FPGAs](#)
- [Introduction to Simulation Acceleration Cosimulation](#)
- [Moving Glue Logic to Improve Performance](#)
- [Support for Final Block in Hardware](#)
- [Moving dumpvars/dumpports to the DUT Automatically](#)
- [Miscellaneous Features](#)

Interface Appearing in Both SW and HW

Interface appearing in both SW and HW. See the following example.

Example

```
interface ift;
endinterface
module dut(ift if_dut);
endmodule
module tb(ift if_tb);
endmodule
module top;
  ift m_if;
  dut dut(m_if);
  tb tb(m_if);
endmodule
```

Moving a HW Module/Interface/Task/Function/Block to Execute in SW

See the following sections for moving a HW module or HW interface to execute in SW while maintaining scope of the intended HW:

- [Using Compile-Time Options](#)
- [Using an Embedded Source Code SystemVerilog Attribute, Comment Pragma, or UTF Command](#)
- [Moving Initial Blocks From HW to SW Automatically](#)

Using Compile-Time Options

It is common to bind verification SystemVerilog interfaces or modules to HW instances or to instantiate an interface in the top-level HW module that has access through a virtual interface from SW, which executes by default in HW.

With this capability, it is possible to move all instances of the specified interface or module to execute in SW.

It might also be helpful to move any HW module in the DUT to execute in SW if it is nonsynthesizable. Precautions should be taken if clock or data races occur.

Example

- For [V2VX Mode Compilation and Runtime Options](#), the following needs to be added to the VCS compile-time options:

```
-Xhdl_cosim_etb <module name>
```

Must be specified for each module or interface that is moved from HW to SW.

- For [V2Z Mode Compilation and Runtime Options](#), the following needs to be added to the UTF file:

```
simxl_move_to_tb -module {<module name list>}
```

All modules specified in the list above executes in SW, using their HW context.

Using an Embedded Source Code SystemVerilog Attribute, Comment Pragma, or UTF Command

Tasks, functions, or blocks implemented in modules or interfaces that are in HW can be moved to execute in SW.

Moving tasks, functions, or blocks from HW to SW is useful for a couple of reasons:

- Usage of code that is not synthesizable or not supported by the Behavioral Compiler
- Code that would be easier to debug executing in SW
- Usage of code that hinders on compile time and frequency (zTime results).

Example

```
module hwmod;
// synopsys simxl_move_tf_to_tb
task task1;
// task contents
endtask
endmodule
```

```
module hwmod;
(* simxl_move_tf_to_tb *)
task task1;
// task contents
endtask
endmodule
```

```
// UTF command example
simxl_move_tf_to_tb -target hwmod.task1

initial begin
  $sformat(inst, "MTK_%m.LOG");
  FID = $fopen(inst, "w");
end

always @ (posedge clk)
  if (<cond>)
    $fdisplay(FID, "VALUE OF DATA IS %h\n", DATA);
```

The example here uses constructs like `$sformat/%m` and other constructs to construct the file name (currently not supported in HW). Adding a pragma before the initial block moves that block to execute in the SW side. File handle can still be used in HW for doing `$fdisplay/$fwrite`. The initial block is moved to HW by adding the attribute `(* simxl_move_blk_to_tb *)`. See the following example.

```
(* simxl_move_blk_to_tb *)
initial begin
  $sformat(inst, "MTK_%m.LOG");
  FID = $fopen(inst, "w");
end
```

To move tasks or function from HW to SW when in the V2VX mode, use the following compile-time option:

```
-simxl=move_tf_to_tb+<module-name>+<task-function-name>
```

Where,

- `<module-name>`: Specifies the name of the module or interface where the task resides in.
- `<task-function-name>`: Specifies the name of the task/function that its execution is being moved from HW to SW.

Moving Initial Blocks From HW to SW Automatically

In a HW module if initial blocks are not synthesizable on ZeBu, you can move specific initial blocks by using the `(* simxl_move_to_tb *)` pragma. However, this method is cumbersome for moving multiple initial blocks.

Simulation Acceleration enables you to automatically move initial blocks, which are not synthesizable on ZeBu, to SW and simulate them in VCS.

For more information, see the following subsections:

- [Enable Automatic Movement of Initial Blocks from HW to SW](#)
- [Movement Considerations](#)

- [Limitations](#)
- [Example](#)

Enable Automatic Movement of Initial Blocks from HW to SW

To move initial blocks to SW side automatically, specify the following VCS elaboration compile-time option:

```
-simx1=move2tb_zebuinit
```

See [Example](#).

Movement Considerations

The following considerations:

- As Simulation Acceleration provides multiple methods for moving initial blocks, you can use all methods. The movement priority is as follows, starting with the highest precedence:
 - Specifying initial block with `(* simx1_not_move_to_tb *)`
 - Specifying initial block with `(* simx1_move_to_tb *)`
 - Using `move2tb_zebuinit` and `move2tb_allinit`

Therefore, even if you use the `move2tb_zebuinit` VCS elaboration compile-time option, initial blocks specified with the `(* simx1_not_move_to_tb *)` pragma are not moved.

- **Performance:** Any initial block that contains an effective `#delay`; that is, the clock delay is enabled in that module by the `clock_delay` UTF command, is not moved. This ensures that high frequency drivers, such as clock generators, are not moved.

Initial blocks that do not have `clock_delay` UTF command enabled, can still be moved to the SW side. Such blocks have no delay and therefore the `#delay` is replaced with `#0`.

Note:

The RTL clock in the transactors cannot be replicated.

See [Example](#).

- **Race Condition:** For both structural and procedural writes, all initial assignments in SW occur before HW events. To avoid race conditions, Simulation Acceleration stops the driver clock in the HW side at the start of cosimulation and restarts the driver clock after all initial values are sent from the SW to the HW.

Limitations

You cannot move an initial block with any task or function call for import tasks in DUT.

Example

This example shows the following:

- Effect of specifying the `-simxl=move2tb_zebuinit` VCS elaboration compile-time option
- Impact of specifying `#delay` in the UTF file

Suppose you have the following file, `test.v`:

```
// SW Module
module testbench ();
  dut top();
endmodule

// HW Module
module dut();
  reg a, b;
  mid mid(a, b);

  initial a = 1; // (1)
  initial #1 b = 1; // (2)

endmodule
...
```

When you compile `test.v` with `-simxl=move2tb_zebuinit`, the initial block (1) and initial block (2) are simulated on the SW side.

To understand the impact of `#delay`, suppose the UTF file contains the `clock_delay` command for the `dut` module, as shown in the following:

```
// dut.utf
...
clock_delay -module dut
```

After compiling `test.v`, only initial block (1) is moved. This is because initial blocks with `#delay` are not moved.

Note that the `#delay` in (2) is ignored in HW because no `clock_delay` is specified. That is, `reg b` is set to 1 at `time=0`. This behavior should be kept after moving. Therefore, Simulation Acceleration moves the following two blocks to SW side.

```
initial a = 1; // (1)
initial #0 b = 1; // (2), #1 is replaced by #0
```

SVA Assertions are Supported in SW and HW That Is Not Mapped to IF FPGAs

SVA assertions are supported in SW and in HW that is not mapped to IF FPGAs . Use the following Simulation Acceleration VCS runtime options for special control on the HW side. Use regular VCS features for assertions executing on SW side.

- Enable HW SVAs:

```
-simxl=enable_dut_sva
```

- Disable specific HW SVAs:

```
-simxl=disable_dut_sva, assertFile=<filename>
```

Where, <filename> contains a list of assertion names to be disabled.

- Disable all assertion in DUT:

```
-simxl=disable_assertion_in_dut
```

VHDL Support and Limitations

This topic contains the following sections:

- [VHDL Support](#)
- [Limitations](#)

VHDL Support

Simulation Acceleration supports SystemVerilog/VHDL mixed designs with VHDL. The support includes:

- VHDL hierarchy in HW:
 - VHDL constructs as supported by ZeBu.
 - Target signal can be port, variable or signal.
 - Selects (bit and part-select) are supported.
- VHDL hierarchy in SW.
- SystemVerilog Related:
 - SystemVerilog hierarchy under VHDL in HW.
 - Any level of SystemVerilog/VHDL nesting in HW.
 - Any level of SystemVerilog/VHDL nesting in the testbench.

- XMR Related:
 - Cross domain XMRs from SystemVerilog SW to VHDL hierarchy in HW.
 - Cross domain XMRs from SystemVerilog SW to SystemVerilog hierarchy under VHDL in HW.
 - XMR read in SW from signal in HW VHDL hierarchy.
 - XMR read in SW from signal in HW SystemVerilog hierarchy under VHDL.
 - XMR force from SW into HW SystemVerilog hierarchy under VHDL.
 - XMR write to HW Verilog hierarchy under VHDL, excluding Verilog ports.
 - XMR through VHDL block statement.
 - Case insensitive specification of names in VHDL hierarchy.
 - Direct XMR for read, force, and write (no generates).
 - Direct XMR for read, force, and write (with VLOG/VHDL generates).
 - Read XMR in `$hdl_xmr` (with `$hw_read`).
 - XMR in `$hdl_xmr_force/$hdl_xmr_release` (with `$hw_force`).
 - Write XMR in `$hdl_xmr` (with `$hw_write`), excluding Verilog ports.
 - Complete shadow hierarchy of MX DUT hierarchy/XMR from VLOG DUT to VLOG TB (DUT2TB)/VLOG embedded testbench under VHDL.
- UCLI Related:
 - UCLI read of VHDL DUT signal from TB
 - UCLI force of VHDL DUT signal from TB
 - UCLI read of VLOG DUT signal from TB (under VHDL)
 - UCLI force of VLOG DUT signal from TB (under VHDL)
- Case insensitive specification of names in the VHDL hierarchy.
- Ability to move a SystemVerilog HW module/instance to SW under VHDL in HW.
- Support for both V2VX and V2Z flows.
- VHDL is supported in testbench hierarchy.
- Support for VLOG embedded testbench under VHDL.

Limitations

The following limitations exist with VHDL support in Simulation Acceleration:

- VHDL in DUT Hierarchy
 - HW top module should be a SystemVerilog module. It cannot be a VHDL module.
 - SW hierarchy above HW top should be Verilog only. It cannot have any VHDL.
- XMRs to VHDL (read/force) can be either direct XMR or specified through `$hdl_xmr` or `$hdl_xmr_force` system calls. However, XMRs to SystemVerilog under VHDL can be specified using direct XMR.
- Write XMRs to VHDL (writing a HW signal from SW) is not supported.
- XMR to complex type, such as records and arrays, is not supported.
- `$hdl_xmr/$hdl_xmr_force` should be accompanied by `$hw_read` and `$hw_force`. This is because `$hdl_xmr/$hdl_xmr_force` can have strings whose value is available only at runtime. However, if the `$hdl_xmr/$hdl_xmr_force` have constant string literals, then use the following compile-time option to let the tool parse the string literal:

```
-simx1=parse_hdl_xmr
```

This avoids the need for `$hw_read` and `$hw_force`.

This functionality might not work for some string literals. For example, it does not work if the string has a `/` instead of a `.` as an XMR separator.

- VHDL architecture cannot be specified as a module that can moved from HW to SE. This is ignored with a warning.
- VHDL cannot be specified as an embedded testbench.
- Incremental compile with VHDL is not supported. However, to force incremental compile, you can specify `-simx1=incr_comp_mx`. This assumes that neither DUT nor communication has changed. This switch has no effect in the first (regular) compile. However, the switch causes the following in subsequent (incremental) compiles:
 - Simulation Acceleration assumes that DUT has not changed.
 - Simulation Acceleration assumes that communication has not changed.
 - Simulation Acceleration does not perform any checks to ascertain the above. You must ensure that it is safe to generate simv from TB hierarchy.
 - Simulation Acceleration does not invoke UFE/Simon or ZeBu backend. simv is generated from TB hierarchy.
- UCLI write of VHDL DUT signal from TB is not supported.

- Task Function Calls:
 - Task Function calls from VLOG TB to VLOG DUT (under VHDL)
 - Task Function calls from VLOG TB to VHDL DUT
 - Task Function calls from VHDL DUT to VLOG TB
 - Task Function calls from VLOG DUT to VLOG TB

Moving Glue Logic to Improve Performance

Glue logic is code in the testbench that connects multiple DUT ports. To increase performance, you might want to run the glue logic in the hardware, instead of the testbench. For example, suppose there is a clock-generator module, which is also marked as DUT. The clock output of the module comes to the testbench and then the clock is inverted and fed to the actual DUT module. In such a case, for performance enhancement, you can move the inverter to the hardware.

Movement of glue logic is always from the testbench to the hardware.

For more information, see the following subsections:

- [Types of Glue logic Supported](#)
- [Moving Glue Logic](#)

Types of Glue logic Supported

Simulation Acceleration supports the cont-assign glue logic. Both cont-assign of simple nodes and cont-assign of XMRs are supported. However, support is provided only for one-level XMRs (target XMRs should be in the direct child of the module where cont-assign is located). XMRs, such as 'TB.clkgen.mid.bot.a', are not supported.

Error messages are reported if you attempt to move the following types of glue logic:

- Primitive gates, such as AND
- Always/Always_comb (sometimes Always_ff/Always_latch)
- User-defined primitives (UDP)

Moving Glue Logic

To move glue logic, perform the following:

1. In the testbench, identify the construct that you want to move.
2. Update the testbench code with the following in-code pragma:

```
(* simxl_move_to_dut *)
```


In the following code snippet, the cont-assign `assign clk_n = ~clk` is identified for movement and consequently simulated in the hardware.

```
// SW Modules
module TB;
wire clk, clk_n;

(* simxl_move_to_dut *) // pragma for moving
assign clk_n = ~clk; // logic to be moved

clkgen clkgen(clk);
dut dut(clk_n);
initial begin
    #10 $finish();
end
endmodule
```

1. Specify the following compile option to see status of the movement:

```
-simxl=diag
```

SIMXL-DIAG messages and other diagnostic information is printed on the screen. See the SIMXL-DIAG message to see if the movement is successful. A sample message is shown as follows:

```
SIMXL-DIAG: Cont-assign ' (* simxl_move_to_dut = 1 *) assign clk_n =
(~clk); ' has been inserted into waitlist.
SIMXL-DIAG: ' (* simxl_moved_to_dut = 1 *) assign clk_n = (~clk); ' has
been moved to tb shadow.
```

1. Open the translog report located in the current working directory to view the improved performance after moving glue logic. You can generate the translog report by using the following runtime option:

```
-simxl=translog
```

For information on translog options, see Debug Options for Communication Mechanisms.

Support for Final Block in Hardware

Simulation Acceleration implements a final block by moving the block from the hardware side to the software side. The following are supported in the final block:

- Access to SW nets and variables
- Call to SW DPI or function/tasks
- Access to HW registers

You can enable this support by setting the following compile option:

```
-simxl=move2tb_final
```

Limitation

Call to HW tasks or functions is not supported.

Support for Concurrent Assertion in Hardware

Simulation Acceleration provides `task/$display` task support in HW hierarchy for Concurrent Assertion failure. The following cases are supported:

- Combination of XMR to task without parameters and display task
- Combination of XMR to task with constant and non-constant parameters and display task
- Combination of a mixture of statements of display and XMR tasks

For more information, see the following subsections:

- [Enabling Concurrent Assertion in Hardware](#)
- [Example](#)
- [Limitations](#)

Enabling Concurrent Assertion in Hardware

To enable `task/$display` in HW modules (trigger execution of the `task/$display` on TB at compile time), set the following compile-time option:

```
-simxl=assert
```

To enable `task/$display` in HW modules (trigger execution of the `task/$display` on TB at runtime), set the following runtime option:

```
-simxl=enable_dut_sva
```

Example

This example shows Concurrent Assertion in SystemVerilog.

```
# Concurrent Assertion
module dut;
.....
task mytask;
$display("Hello World!!");
endtask
...
always @(negedge clk)
```

```
B1: begin
S1:assert property (p1(select,gnt_1,gnt_2))
    $display("Assertion P1 Passed");
else begin
    mytask();
end

S2:assume property (p1(select,gnt_1,gnt_2))
    $display("Assume P1 Passed");
else
    mytask();
end
.....
endmodule
```

Limitations

For the following limitation, the Simulation Acceleration tool reports error messages:

- Assertion fail cannot have statement other than task calls or \$display system tasks:

Moving dumpvars/dumpports to the DUT Automatically

In the ZeBu flow, signals that need to be written out are usually specified in a separate file that contains the \$dumpvars or \$dumpports specifications. However, in the Simulation Acceleration flow, the module can be marked as a TB module. Such modules are ignored by ZeBu, if you have not marked the module as a DUT top in the UTF file. To avoid this overhead, the Simulation Acceleration flow supports specifying the \$dumpvars and \$dumpports system tasks in the testbench.

The Simulation Acceleration tool identifies \$dumpvars/\$dumpports system task statements in an initial block of a TB module with a target of DUT hierarchy. These \$dumpvars/\$dumpports system task calls are moved to the DUT. The following cases are considered:

- If the \$dumpvars/\$dumpports is in a standalone module, the partition of the module is changed from TB to DUT. Therefore, the whole module is moved to the DUT side.
- If the \$dumpvars/\$dumpports is in a dedicated initial block, but the parent module has other constructs, the Simulation Acceleration tool moves the dedicated initial block with \$dumpvars/\$dumpports to a newly created DUT module, which is then used to save signals. Otherwise, if the \$dumpvars/\$dumpports is not in a dedicated initial block, the movement does not occur.

The movement is performed only in the following cases:

- When `$dumpvars/$dumpports` system task statements are in an initial block of a TB module with a target of DUT hierarchy
- When no Simulation Acceleration Not Yet Implemented (NYI) messages are reported.

For more information, see the following subsections:

- [Example 1: Dedicated Initial Block](#)
- [Example 2: Non-Dedicated Initial Block](#)
- [Limitations](#)

Example 1: Dedicated Initial Block

This example shows a `$dumpvars` specification in a TB module. The specification is in a dedicated initial block.

```
module TB
...
not n1(out, in); // other constructs
initial in = 1'b1; // other constructs
initial begin // dedicated for dumping
    $dumpvars(0, DUT);
end
```

Here, the `$dumpvars` specification has the following arguments:

- 0: Specifies that all signals need to be written out
- DUT: Specifies the module name or instance name

Only the dedicated initial block is moved instead of the whole module to the DUT side. The Simulation Acceleration tool creates destination module for moving is a special module called `hwcosim_dumpvar_mod`.

Example 2: Non-Dedicated Initial Block

This example shows a `$dumpvars` specification in a TB module. The specification is in a non-dedicated initial block.

```
module TB
...
initial begin // not a dedicated initial block
    a = 1'b0;
    $display(a);
    $dumpvars(0, DUT);
end
```

Limitations

The following specification is not supported. If the `$dumpvars/$dumpports` refers both TB and DUT instances, the `$dumpvars` are not moved because the movement might invalidate the writing of for TB instances. The Simulation Acceleration tool reports an error message.

```
module TB
...
not nl(out, in); // other constructs
initial in = 1'b1; // other constructs
initial begin // dedicated for dumping
$dumpvars(0, sub.a.b, tb.a);
end
```

Here, `$dumpvars` has the following arguments:

- 0: Specifies that all signals need to be written out
- `sub.a.b`: Hierarchy under the hardware
- `tb.a`: Hierarchy under the software

In this case, only the signals on TB side are written out.

Miscellaneous Features

- Packages shared by HW and SW to enable common data types
- Packages used only by HW with supported HW constructs
- HW subroutines in an interface can call other HW subroutines in a different interface
- There is no limit on the nested calls of subroutines between HW and SW. For example, SW calls a HW subroutine, that calls a SW class method, and so on
- VHDL DUT wrapped in a SystemVerilog module is supported in HW

In Simulation Acceleration, the access limitations that do not include signals or subroutines are as follows:

- [Other Limitations That Do Not Include Signals or Subroutines](#)
- [Temporary Known Simulation Acceleration Limitations](#)

Other Limitations That Do Not Include Signals or Subroutines

See the following Simulation Acceleration limitations:

- Accessing SystemVerilog events between SW and HW

```
interface ift;
event ev;
task task1(output o1);
begin
->ev;
end
endtask
function integer calculate();
return(0);
endfunction
endinterface
module HW();
ift m_if();
endmodule
module SW();
logic a;
initial
begin
@(HW.m_if.ev); // Illegal
$display("A: %d, CALC: %d", a, HW.m_if.calculate());
end
endmodule
```

- Accessing SystemVerilog chandle variable between SW and HW

chandle is not supported in HW.

```
chandle variable_name;
```

- SW disabling a named block in HW

```
interface ift;
always @(posedge clk)
begin: myblock
...
end
endinterface
module HW();
ift m_if();
endmodule
module SW();
initial
begin
disable HW.m_if.myblock; // Illegal
end
endmodule
```

- **Accessing SystemVerilog dynamic types between SW and HW**

Dynamic types are not supported in HW.

```
interface ift;
byte amem [int]; // Illegal
task task1(output o1);
begin
o1 = 1;
end
endtask
function integer calculate();
return(0);
endfunction
endinterface
module HW();
ift m_if();
endmodule
module SW();
logic a;
initial
begin
a = HW.m_if.amem[0][0] // Illegal HW.m_if.task1(a);
$display("A: %d, CALC: %d", a, HW.m_if.calculate());
end
endmodule
```

- **Accessing SW elements from HW that reside in a SystemVerilog package** (including the default \$unit package) except for data types that are support in HW and class handles used for calling import methods that reside in a class object.
- **VCS system task \$set_toggle_region is not supported**

```
// None of the VCS system tasks below are supported.
initial begin
$read_lib_saif("inputFile");
$set_gate_level_monitoring("on");
$set_toggle_region(Scope);
$toggle_start;
$toggle_stop;
$toggle_report("outputFile", timeUnit, Scope);
end
```

- **No support for \$dumps in SW referencing HW hierarchies**
- **No support for same interface in both HW and SW**, where HW instantiation is being referred by virtual interface.

```
interface ift;

task task1(output o1);

begin
```

```
        o1 = 1;
    end

    endtask

    function integer calculate();
        return(0);
    endfunction

    endinterface

    module HW();
        ift m_if();
    endmodule

    module SW();
        logic a;

        virtual interface vif;

        ift m_if_sw();

        initial

        begin

            vif = HW.m_if;

            $display("A: %d, CALC: %d", a, vif.calculate());

        end

    endmodule
```

- **\$stop is not supported**
- **Usage of Clock Delay Ports in Simulation Acceleration flow is not recommended.** Clock Delay Ports timescale by default is 1 ps/1 ps and cannot be changed. This does not work well with other timescales.
- **tran, tranif, or rtranif in SW connected to the DUT boundary port, which acts as a true bidirect:** Simulation Acceleration does not support tran, tranif, or rtranif in

SW connected to the DUT boundary port, which has drivers in both SW and HW (true bidirect).

- Simulation Acceleration ignores a signal that is initialized to 1 in an initial block at zero time in a transactor.

Temporary Known Simulation Acceleration Limitations

Following are known limitations with the current release of Simulation Acceleration:

- To enable debug for Simulation Acceleration, use the following UTF command:

```
debug -verdi_db true
```

Do not use the `debug -all true` UTF command to enable debug for Simulation Acceleration.

- In testbench, if there are some events happening at the same time as `$finish` in a different thread, there could be difference in the order of execution of the threads between simulation and acceleration. In this case, a small delay to ensure `$finish` executes at the end.
- In a transactor, at zero time when a signal is initialized to 1 in an initial block, simulation considers it as a posedge where as SimXL ignores it.
- In testbench, if there are some events happening at the same time as `$finish` in a different thread, there could be difference in the order of execution of the threads between simulation and acceleration. In this case, you might add a small delay to ensure `$finish` executes at the end.

2

Supported Syntax in Simulation Acceleration

For more information, see the following topics:

- [Supported Syntax for Using Signals as the Communication Mechanism](#)
- [Signal Communication Using Primitive Data Types](#)
- [Supported Syntax for Using Subroutines as the Communication Mechanism](#)
- [Supported and Unsupported Behavioral Compiler Syntax](#)

Supported Syntax for Using Signals as the Communication Mechanism

In Simulation Acceleration, the elements supported for signal-level communication between HW and SW are described in the following subsections:

- [Signal Communication Using Primitive Data Types](#)
- [Signal Communication Using Aggregate Data Types](#)
- [Signal Communication Using Interface Ports](#)
- [Signal Communication Using XMRs](#)
- [Optimizing Performance for Clock XMRs From SW to HW \(clkopt\)](#)
- [Optimizing Performance for XMRs With Force and Release Statements](#)
- [SW Accessing HW Memory Using \\$readmem and \\$writemem System Tasks](#)
- [Accessing and Forcing HW Signals From Within SW Using \\$hdl_xmr and \\$hdl_xmr_force](#)
- [Signal Communication Through External Subroutines](#)
- [Strength Analysis in Simulation Acceleration](#)
- [Limitations With Signal-Based Communication](#)

Signal Communication Using Primitive Data Types

Signal communication using the following primitive data types for signals:

- wire, wand, wor, tri, and so on.
- reg
- logic
- enum
- byte
- integer
- int
- shortint
- longint

Signed modifiers can be applied to any of these data types.

Example

```
module HW(input wire clock, input logic reset);
endmodule
interface ift;
logic clock;
assign HW.clock = clock;
endinterface
module SW;
ift m_if();
endmodule
```

Signal Communication Using Aggregate Data Types

Signal communication using the following aggregate data types:

- Packed arrays
- Unpacked arrays
- Multidimensional arrays
- Packed structs
- Unions

Any composition of these data types is supported in Simulation Acceleration.

Example

```
typedef struct packed{
integer source;
integer destination;
} descriptor;
module HW(input wire clock, input logic reset, input wire descriptor
d[100]);
endmodule
interface ift;
descriptor m_d[100];
genvar i;
for (i=0; i<100; i++)
assign HW.d[i] = m_d[i];
endinterface module SW;
ift m_if();
endmodule
```

Signal Communication Using Interface Ports

Signal communication using interface ports:

- Interface ports containing signals with any of the types stated in the previous sections
- Modports

Example

```
interface ift;
endinterface
module dut(ift if_dut);
endmodule
module tb(ift if_tb);
endmodule
module top;
ift m_if;
dut dut(m_if);
tb tb(m_if);
endmodule
```

Signal Communication Using Real Types

The Simulation Acceleration flow supports:

- Ports with `real` type at the boundary between the testbench and HW
- Predefined `nettype` with either of the following types: `current_r`, `voltage_r`, `wrealldriver`, `wreal4state`, `wrealavg`, `wrealmax`, `wrealmin`, `wrealsum`.
- An Embedded testbench (ETB) module with `real` port or predefined `nettype` with `real` type

Use Model

To enable this support, specify the following `vcs` switch:

```
-simx1=enable_realport
```

Limitations

The following is not supported:

1. Inout port with `nettype`
2. Cross-boundary `xmr` to `real` type variable or `nettype`
3. Task/function/DPI calls between testbench and HW, which have the following argument type:
 - a. `nettype` or array of `nettype`
 - b. struct containing `real` or array of `real`
4. Cross-boundary force/release on the `real` type or `nettype` node

Signal Communication Using XMRs

XMR support is enabled by default. The following table shows the supported elements for signal communication using XMRs.

Table 1 Supported Elements for Signal Communication Using XMRs

Supported Elements	Example
Signal communication using XMRs with structural code XMRs between SW and HW using continuous assignment and port connections.	<pre>module dut(input a); endmodule module tb; assign dut.a = 1'b1; endmodule</pre>

Table 1 *Supported Elements for Signal Communication Using XMRs (Continued)*

Supported Elements	Example
Signal communication using XMRs with behavioral code XMRs from SW to HW for reading HW signal values of any type. Note that behavioral reads of data types mapped to ZeBu memories are also supported in Simulation Acceleration.	<pre> module dut(input a); reg b; endmodule module tb; reg c; initial c = dut.b; endmodule </pre>
Signal communication using XMRs with behavioral code XMRs from SW to HW for writing to HW signals of any data type. Note that behavioral writes of data types mapped to ZeBu memories are also supported in Simulation Acceleration.	<pre> module dut(input a); reg b; endmodule module tb; reg c; initial dut.b = c; endmodule </pre>
Signal communication using XMRs with behavioral code XMRs from SW to HW for applying force or release to HW signals of any data types.	<pre> module dut(input a); reg b; endmodule module tb; reg c; initial force dut.b = c; endmodule </pre>

Table 1 Supported Elements for Signal Communication Using XMRs (Continued)

Supported Elements	Example
Signal communication using XMRs that require a named block as a segment in the hierarchical access.	<pre> module sub(); reg b; endmodule module dut(input a); genvar i; for (i=0; i<10; i++) begin : NESTED sub sub(); end endmodule module tb; reg c; initial dut.NESTED[1].sub.b = c; endmodule </pre>

Optimizing Performance for Clock XMRs From SW to HW (clkopt)

Clocks are high frequency signals that cause frequent synchronization between SW and HW when used in SW to reference HW clocks through XMRs, slowing down performance. When using HW clock references in a repeat statement, there is only a need to synchronize between SW and HW only when the specified number of clock iterations complete. To enable an optimization such that there are not frequent synchronizations, use the following VCS elaboration compile-time option:

```
-simx1=clkopt
```

```
// Example of repeat statement reference HW clock in testbench code
repeat (100) @(posedge vif.clock);
```

Support for iff Construct with clkopt

The event control to be optimized by `clkopt` can have the `iff` construct. In this case, `clkopt` moves the entire (clock and condition) conditional event control to HW.

The reference to clock and condition should be via a virtual interface where the same virtual interface is used in both the clock and condition.

```
@(posedge vif.clock iff vif.rst==1'b1);
```

Enhancements to clkopt Optimization

The `clkopt` optimization is enhanced to support:

- body of repeat statement and
- an if-else block, with one having posedge and other having negedge clock.

The following code snippet shows the new support:

```
task automatic wait_n_edges(int rep, bit pos_edge);
  repeat(rep) begin
    if (pos_edge) @(posedge tb.d1.pl_s_clk);
    else @(negedge tb.d1.pl_s_clk);
  end
endtask
```

Optimizing Performance for XMRs With Force and Release Statements

This optimization avoids expensive SW to HW accesses needed for force assignments by moving the execution of the force/release to HW.

For example, suppose the testbench has a `force` statement as shown in the following example:

```
force tb.dut.xmr.to.lhs = tb.dut.xmr.to.rhs;
```

Similarly, a `release` statement can exist in the testbench.

The optimization moves the execution of the assignment above to HW because both LHS and RHS are XMRs to the HW.

To enable this optimization, use the following VCS elaboration compile-time option:

```
-simxl=forceopt
```

SW Accessing HW Memory Using \$readmem and \$writemem System Tasks

SW accessing HW memory using `$readmem` and `$writemem` system tasks is supported in Simulation Acceleration. Moving a HW initial block containing `$readmemh` from HW to SW is supported. The `$readmemh` and `$writemem` system tasks are ignored if placed in HW.

All variations of these tasks are supported:

- `$writememb`, `$writememh`
- `$readmemb`, `$readmemh`

The start and finish address arguments are optional as specified in SystemVerilog LRM.

Example

```
interface ift;
logic [0:255] mem [4096];
endinterface
module dut;
endmodule
module HW;
ift m_if();
dut dut();
endmodule
module SW;
initial
begin
$readmemh("data.hex", HW.m_if.mem);
$writememh("data_copy.hex", HW.m_if.mem);
end
endmodule
```

Accessing and Forcing HW Signals From Within SW Using \$hdl_xmr and \$hdl_xmr_force

The following system tasks are supported in Simulation Acceleration:

- [\\$hdl_xmr: Access HW Signals From Within SW](#)
- [\\$hdl_xmr_force: Force HW Signals From Within SW](#)

\$hdl_xmr: Access HW Signals From Within SW

The syntax of the `hdl_xmr` procedure is as follows:

```
$hdl_xmr("source_object", "destination_object", [verbosity]);
```

Where:

- `source_object`: Specifies a VHDL signal or Verilog register or net using a string type. You can specify an absolute path or a relative path to the object.

Note:

Use an absolute path instead of a relative path, if the source node resides in the VHDL part of the code or if the hierarchical path has a VHDL layer.

- `destination_object`: Specifies a VHDL signal or a Verilog register using a string type. You can specify an absolute path or a relative path to the object.

\$hdl_xmr_force: Force HW Signals From Within SW

You can use `$hdl_xmr_force()` calls to force a value on an existing Verilog or VHDL node. It allows you to force signals, from and at any level of the design hierarchy from within the VHDL architecture or the Verilog module and from any scope from where the standard defined system functions can be called. The `$hdl_xmr_force()` procedure works similar to the Verilog `force` command. A status is returned from this function, which can be used to control the flow of the code.

The syntax is as follows:

```
$hdl_xmr_force(destination_object, value);
```

Where:

- `destination_object`: Specifies the hierarchical path as a string. A full hierarchical path (or relative path with reference to the calling block) to an existing VHDL signal or Verilog register/net.
- `value`: Specifies the value to which the `destination_object` is to be forced. The value should be specified within double quotes as a constant string type. The specified value must be appropriate for the type and be specified in valid syntax for any radix, as supported by the language.

Signal Communication Through External Subroutines

Signal communication through external subroutines, such as PLI, VPI, and UCLI, is supported in Simulation Acceleration.

Declare signals in software to ensure the type of access is known at compile time. Use the `$hw_read`, `$hw_write`, and `$hw_force` VCS built-in system tasks. The `$hw_read` system task call takes an additional optional string literal argument called “passive”, which causes the signal to be communicated passively, instead actively.

By default, XMR communication is active. This indicates that the HW sends the value of the signal to the SW whenever the value changes. However, in passive communication, the HW sends the signal to the SW along with an existing active XMR communication or a sync from SW to HW. Setting XMR communication to passive is less resource intensive.

Passive communication can be done if the usage of the signal of the XMR is in procedural context. If the usage is in structural context and the XMR communication is set to passive, the Simulation Acceleration tool reports an error.

Simulation Acceleration does not report an error for the following uses of XMR, when passive is specified:

- `$display(dut.xmr);`
- `initial lhs=dut.xmr;`
- `always @(posedge clk) lhs=dut.xmr;`

Simulation Acceleration reports an error for the following uses of XMR, when passive is specified:

- `assign lhs=dut.xmr;`
- `dut dut(dut.xmr); //usage in port`

Example 1

```
initial begin
$hw_read(top.dut.w);
$hw_write(top.dut.u);
$hw_force(top.dut.v);
end
```

Example 2

This example sets the XMR communication to passive.

```
$hw_read(dut.xmr, "passive");
```

Strength Analysis in Simulation Acceleration

Strength analysis is used to check the strength of a wire in the HW. Strength is propagated when signals are passed through port connections from the SW to HW.

You can check the strength propagation for boundary ports and XMRs. When using Simulation Acceleration, you can set the strength of a net in the SW and consequently, check the strength of the net in the HW using the following ways:

- **Cont-assign:** See the following example specification:

```
assign (weak1, weak0) net1 = a & b;
```

- **Primitive-gate:** See the following example specification:

```
bufif1 (pull1, pull0) net2 = c & d;
```

- **MOS-gates:** See the following example specification:

```
rnmos(net3, in, ctrl);
```

Note:

Strength analysis for triregs is not supported.

Enabling Strength Analysis

Set the following compile time option to enable strength propagation for boundary ports and XMRs:

```
-simx1=enable_strength
```

Viewing Strength Results in HW

At runtime, you can check the strength passed from SW to the HW (DUT) by viewing details in translog. Use the following runtime option:

```
-simx1=translog
```

For information on translog options, see the Debug Options for Communication Mechanisms.

The strength is visible in the translog report, as shown in the following snippet:

```
SimTime: 0
[X] [Write] top.sd.d._vcs_hwcosim__tb2dut_init_wait [0:0] 'd0
[X] [Write] top.sd.d._vcs_hwcosim__tb2dut_init_wait [0:0] 'd1
[P] [SwHw] top.sd.d.in1 [0:0] 'bx, strength1 = 6, strength2 = 3
[P] [SwHw] top.sd.d.in2 [0:0] 'bx, strength1 = 6, strength2 = 7
[P] [SwHw] top.sd.d.in2 [1:1] 'bx, strength1 = 5, strength2 = 6
[P] [SwHw] top.sd.d.in2 [2:2] 'bx, strength1 = 6, strength2 = 5
[P] [SwHw] top.sd.d.clk [0:0] 'd0
SimTime: 2
[P] [SwHw] top.sd.d.in2 [0:0] 'd0, strength1 = 7, strength2 = 7
[P] [SwHw] top.sd.d.in1 [0:0] 'd0, strength1 = 3, strength2 = 3
[P] [SwHw] top.sd.d.in2 [2:2] 'd0, strength1 = 5, strength2 = 5
[P] [SwHw] top.sd.d.in2 [1:1] 'd0, strength1 = 6, strength2 = 6
```

Limitations With Signal-Based Communication

The following elements are not supported in signal-level communication between HW and SW.

- Signal communication with XMR using wreal, real, and shortreal data types.

An example of code that is not supported is as follows:

```
module dut();
real a;
```

```
endmodule
module tb;
assign dut.a = 1.05;
endmodule
```

- Signal communication with port or XMR using time data type.

An example of code that is not supported is as follows:

```
module dut(input time a);
endmodule
module tb;
assign dut.a = $time;
endmodule
```

- Signal communication with port or XMR using tagged unions.

An example of code that is not supported is as follows:

```
typedef union tagged {
void Invalid;
int Valid;
} VInt;
module dut(input time a);
VInt vi1, vi2;
endmodule module tb;
assign dut.v1 = tagged Valid (23+34);
endmodule
```

- Signal communication with signals in an interface have the following limitations :
 - Should not be a signal declared within clocking blocks
 - The interface should not have subroutines that are called from HW
 - The interface should not have signals driven simultaneously by both SW and HW
 - The interface should not be passed as port argument using a XMR

An example of code that is not supported is as follows:

```
interface ift;
clocking cb @(posedge clock);
endclocking
task a;
endtask endinterface
module dut(ift if_dut);
endmodule
module tb(ift if_tb);
endmodule
module top;
ift m_if;
dut dut(m_if);
```

```
tb tb(m_if);  
endmodule
```

- Signal communication with port or XMR using signal references containing segments using named blocks not from generate statements.

An example of code that is not supported is as follows:

```
module top_legal;  
  initial begin  
    for (int i=0; i<3; i++) begin: illegal //automatic int loop3=0;  
      for (int j=0; j<3; j++)  
        begin  
          loop3++;  
          $display(loop3);  
        end  
      end  
    end  
  endmodule  
module tb;  
  initial  
    top_legal.illegal.loop3 = 0; // illegal  
endmodule
```

- Unpacked structs are not supported for signal-based communication.

Supported Syntax for Using Subroutines as the Communication Mechanism

For more information on the supported syntax for using subroutines as the Simulation Acceleration communication mechanism, see the following subsections:

- [Interface Task/Function \(Subroutines\) Calls From SW to HW](#)
- [Class Method Calls From HW Interfaces to SW](#)
- [DPI Subroutine Call From HW Interfaces](#)
- [Streaming Export Methods](#)
- [UVM Messaging Support](#)
- [Import Calling an Export](#)
- [\\$test\\$plusargs, \\$value\\$plusargs Support](#)
- [System Tasks](#)
- [System Task Support for \\$display](#)
- [Enabling Concurrent Export Task/Function Calls](#)

Interface Task/Function (Subroutines) Calls From SW to HW

Interface Task/Function (subroutines) calls from SW to HW:

- Interface subroutines should only consist of syntax supported by the Behavioral Compiler (see [Supported and Unsupported Behavioral Compiler Syntax](#)).
- The same interface instance subroutine cannot be called concurrently from two different places, regardless of whether it is automatic or not (see option to enable calling of concurrent subroutines in [Enabling Concurrent Export Task/Function Calls](#)).
- Access from SW to HW interface subroutine is done either by using a physical hierarchical XMR to the interface or using a virtual interface as a prefix to access the subroutine in the interface.
- Formal arguments of subroutines can be of any of the data types supported for signal communication mentioned in [Supported Syntax for Using Signals as the Communication Mechanism](#).
- To get values back through SW calls to HW tasks, use the output and inout modifiers.
- Ref modifier is not supported.
- Time arguments are supported when *clock_delay* is used. See [Supported and Unsupported Behavioral Compiler Syntax](#).
- This kind of call is often referred to as an "export call".
- Any interface or module that has an export task defined in it is referred to as a transactor.

Example

```
interface ift;
task task1(output o1);
begin
o1 = 1;
end
endtask
function integer calculate();
return(0);
endfunction
endinterface
module HW();
ift m_if();
endmodule
module SW();
logic a;
initial
begin
HW.m_if.task1(a);
```

```
$display("A: %d, CALC: %d", a, HW.m_if.calculate());  
end  
endmodule
```

Class Method Calls From HW Interfaces to SW

Class method calls from HW interfaces to SW:

- A SystemVerilog Class object handle must be passed from SW to HW. HW uses the handle to call class methods implemented in SW.
- The same method can be called concurrently from multiple places.
- Formal arguments of methods can be of any of the data types supported for signal communication mentioned in [Supported Syntax for Using Signals as the Communication Mechanism](#).
- Ref modifier is not supported.
- Time arguments are supported when *clock_delay* is used. See [Supported and Unsupported Behavioral Compiler Syntax](#).
- This type of call is often referred to as an "import call".
- Any interface or module that has an import task call is referred to as a transactor.

Example

```
class C;  
function void observe(input integer i1);  
$display("I: %d", i1);  
endfunction  
endclass  
interface ift;  
bit configure_done = 0;  
C m_handle;  
function void configure(C handle);  
m_handle = handle;  
configure_done = 1;  
endfunction  
initial begin  
wait (configure_done == 1);  
m_handle.observe(4);  
end  
endinterface  
module HW();  
ift m_if();  
endmodule  
module SW();  
C c1 = new();  
initial
```



```
begin  
HW.m_if.configure(c1);  
end  
endmodule
```

DPI Subroutine Call From HW Interfaces

DPI subroutine call from HW interfaces:

- The same subroutine can be called concurrently from multiple places.
- Formal arguments of subroutines can be of any of the data types supported for signal communication mentioned in [Supported Syntax for Using Signals as the Communication Mechanism](#).
- Ref modifier is not supported.

Streaming Export Methods

To enable streaming between SW and HW calls, use the following compile-time option:

```
-simxl=nonblocking_tf.
```

With this option, if a HW task/function has none of the following:

- Return value
- Output argument
- Usage of #-delays

The SW to HW calls of that task/function is made streaming. With streaming, the VCS scheduler does not block and wait for HW to complete execution of the task, resulting in faster runtimes.

You can also use attribute (* simxl_nonblocking=1 *) to make a specific task/function as non-blocking.

UVM Messaging Support

The following UVM report macros and function calls are supported in the HW:

- ``uvm_info`
- ``uvm_warning`
- ``uvm_error`

Chapter 2: Supported Syntax in Simulation Acceleration

Supported Syntax for Using Subroutines as the Communication Mechanism

- ``uvm_fatal`
- `uvm_pkg::uvm_report_enabled`
- `uvm_pkg::uvm_report_info`
- `uvm_pkg::uvm_report_warning`
- `uvm_pkg::uvm_report_error`
- `uvm_pkg::uvm_report_fatal`

The UVM package is deduced as a SW module. Therefore, the calls to these functions are treated as HW to SW function calls.

To facilitate UVM messaging in HW, you can use the `$sformatf/$sformat` system tasks to compose string messages in HW.

Example

```
interface hwmonitor(input a, input b);

initial
  forever
    if (a != b)
      `uvm_error("hwmonitor", "Variable a mismatches to b");
    else
      `uvm_info("hwmonitor", "Variable a matches to b");
endinterface
```

All existing functionality of `uvm_message` is consistent with messages coming from SW or HW. This includes message enumeration and runtime verbosity control using the regular runtime option: `+UVM_VERBOSITY=<verbosity-level>`.

Compile time verbosity is enabled using the following VCS elaboration option:

```
-simx1=UVM_COMPILE_VERBOSITY+<verbosity-level>
```

Simulation Acceleration compiles out the UVM message calls whose verbosity level cannot be met.

If there is no verbosity level provided (for example, `-simx1=UVM_COMPILE_VERBOSITY`), Simulation Acceleration assumes the default value of `UVM_MEDIUM`, which is the same as the runtime `+UVM_VERBOSITY` option.

Example

```
int verbosity = 400;
`uvm_info(id1, "MY MESSAGE1", UVM_LOW)
`uvm_info(id2, "MY MESSAGE2", UVM_MEDIUM)
`uvm_info(id3, "MY MESSAGE3", UVM_HIGH)
`uvm_info(id4, "MY MESSAGE4", verbosity)
`uvm_info(id5, "MY MESSAGE5", UVM_DEBUG)
```

In this example, by default all these UVM reporting macros are compiled into the design. If you provide `-simx1=UVM_COMPILE_VERBOSITY+UVM_MEDIUM`, Simulation Acceleration would compile out the following macros:

```
`uvm_info(id3, "MY MESSAGE3", UVM_HIGH)
`uvm_info(id5, "MY MESSAGE5", UVM_DEBUG)
```

It does not delete the macro ``uvm_info(id4, "MY MESSAGE4", verbosity)` because its verbosity argument is not a compile constant.

The use model of variable verbosity can also be applied to other macros such as `uvm_warning`, `uvm_error` and `uvm_fatal`.

If you provide `-simx1=UVM_COMPILE_VERBOSITY+UVM_NONE`, `uvm_info`, `uvm_warning`, `uvm_error` and `uvm_fatal` with verbosity greater than `UVM_NONE` are removed.

If you provide `-simx1=UVM_COMPILE_VERBOSITY+no`, `uvm_info`, `uvm_warning` and `uvm_error` are removed regardless of their verbosity. However, `uvm_fatal` is retained.

Import Calling an Export

Import calling an export (HW calling a SW method that calls a HW subroutine):

- VCS runtime does not allow an import function to call a time-consuming export task or function. VCS runtime allows an import task to call a time-consuming export task or function.
- When an import task or function calls an export, the import must be a task. Simulation Acceleration adds a wait statement in a wrapper function that activates the export task or function, which causes it to be time-consuming. Therefore, the import must be a task.

Example

```
function void dutTaskOne; // export wrapper on SW

input logic i_x;
input logic i_y;
output logic o_x;
output logic o_y;
automatic bit[1:0]
vcs_hwcosim_packedInputs;
begin
  _vcs_hwcosim_0_isHwDone = 1'b0;
  begin
    _vcs_hwcosim_packedInputs = { >> 1{i_x, i_y}}; // pack inputs
  end
  $$hwcosimRxData(_vcs_hwcosim_idbId[0], _vcs_hwcosim_packedInputs); //
    initiate call to HW; actual call happens in preLER region
  wait ((_vcs_hwcosim_0_isHwDone === 1'b1)); // wait until HW is done
end
```

```
{ >> 1 {o_x, o_y}} = _vcs_hwcosim_0_packedOutputs[33:32]; // unpack
  outputs
end
endfunction
```

\$test\$plusargs, \$value\$plusargs Support

`$test$plusargs` are used to check if a runtime option is being used on the `simv` command line. `$value$plusargs` is used to obtain the value of a runtime option being provided on the `simv` command line.

These system tasks are supported in SW, transactor HW modules and in regular HW modules.

To enable the support in both HW module types, add the following lines to the UTF file:

```
system_tasks -task {$test$plusargs} -enable
system_tasks -task {$value$plusargs} -enable
```

Example

```
integer i1;
always@(posedge clk) begin
  if ($test$plusargs("HELLO"))
    $display("Hello argument found.");
  if ($value$plusargs("TEST=%d", i1))
    $display("value was %d", i1);
end
```

Limitation

`$test$plusargs` or `$value$plusargs` are not supported at time 0. For example, the following scenario is not supported:

```
initial
begin
  if ($test$plusargs("myarg")) $display("hello");
end
```

System Tasks

The following system tasks are supported:

- `$display`, see the [System Task Support for \\$display](#) section.
- `$fdisplay`, together with: `$fdisplayh`, `$fdisplayb`, `$fdisplayo`
- `$write`, together with: `$writeh`, `$writeb`, `$writeo`
- `$fwrite`, together with: `$fwriteh`, `$fwriteb`, `$fwriteo`

- `$fopen`
- `$fclose`
- `$random`, together with: `$urandom`, `$urandom_range`
- `$finish`
- `$hdl_xmr`, see the [\\$hdl_xmr: Access HW Signals From Within SW](#) section for usage.

Usage requires adding the system task that is used in UTF file.

Simulation Acceleration has its own way of implementing system tasks, which results in the module/interface using the system task to become a Transactor processed by VCS instrumentation.

To return to regular ZeBu support for system tasks and disable the Simulation Acceleration implementation, add the following option to VCS compile:

```
-simxl=streaming_display
```

Example

```
// UTF commands to enable system tasks
system_tasks -task "$display" -enable
system_tasks -task "$write" -enable
```

Note:

The file I/O operations are better done on the SW side even though supported for Transactor Code.

System Task Support for `$display`

The following system tasks are supported:

- `$display`
- `$displayh`
- `$displayb`
- `$displayo`

The Simulation Acceleration tool can optimize `$display` by moving the display task to the shadow hierarchy of the corresponding module using a HW trigger. Display optimization takes precedence over `'system_tasks -task {$display} -enable'`. If a `$display` cannot be optimized, it falls back to the `system_tasks` implementation.

For more information on display optimization, see the following subsections:

- [Enabling \\$display Optimization for All HW Modules](#)
- [Enabling \\$display Optimization for Specific HW Modules](#)
- [Enabling \\$display Optimization for Backdoor Communication](#)
- [Pragma Support](#)
- [UVM Report Support](#)
- [Limitations of Display Optimization](#)

Enabling \$display Optimization for All HW Modules

To enable `$display` optimization for all HW modules, specify the following compile-time option:

```
-simxl=displayopt
```

Enabling \$display Optimization for Specific HW Modules

When display optimization is used, there is more overhead because of increased communication between HW and SW. This leads to a reduction in performance. To increase the performance, you can enable `$display` optimization for specific HW modules.

To do this, specify the list of modules in a text file and pass the file using the following compile-time option:

```
-simxl=displayopt+modules=<text-file-containing-list-of-modules>
```

Where, the `<text-file-containing-list-of-modules>` contains a module name in each line of the text file. Hash character is a comment and can be specified at the beginning of the line. In addition, blank lines are ignored.

Example:

```
-simxl=displayopt+modules=modules.txt
```

Enabling \$display Optimization for Backdoor Communication

By default, Simulation Acceleration transactors are used to communicate the arguments of the optimized `$display` from HW to SW. However, by using the following compile-time option, you can enable SW to use backdoor read mechanisms to read the values from the HW:

```
-simxl=displayopt+backdoor
```

This can lead to an improvement in performance.

Pragma Support

You can annotate individual `$display` statements using pragma to provide fine control over `$display` optimizations. Pragma optimizations override module-level `$display` optimizations.

The following pragmas are supported:

Table 2 *Pragma Support for \$display Optimizations*

Pragma	Function
<code>(* simxl_displayopt *)</code> or <code>(* simxl_displayopt = "xtor" *)</code>	Enables <code>displayopt</code> for a specific <code>\$display</code> to use transactor for communication. Note: <code>\$display</code> in loops always use backdoor.
<code>(* simxl_displayopt = "backdoor" *)</code>	Enables <code>displayopt</code> for a specific <code>\$display</code> to use backdoor for communication.
<code>(* simxl_displayopt = "none")</code>	Disables <code>displayopt</code> for a specific <code>\$display</code> .

UVM Report Support

In addition to `$display`, the display optimization supports the following UVM reporting tasks:

- `uvm_report_info`
- `uvm_report_warning`
- `uvm_report_error`
- `uvm_report_fatal`

If you are using UVM macros, such as ``UVM_INFO`, the macro calls the `uvm_report` task inside an if condition that checks for the `uvm_report_enabled()` function.

Example

```
always @(posedge clk)
begin
if (uvm_report_enabled(UVM_HIGH))
    uvm_report_info("ID4",
        $psprintf("wait_for_idle_credits() start"),
        UVM_HIGH);
end
```

Limitations of Display Optimization

Display optimization does not support the following:

- **Argument type:** Class, Struct, Interface, enum, memory, function xmr
- **Format specifier type:** %p is not supported
- **Blocks:** The following are not supported:

- System task in combinational blocks, as follows:

```
always_comb
always @(*)
always @(in1 || in2...)
```

- System task in function which used in combinational statements, as follows:

```
assign x = func(a,b);
function byte func(int x, int y);
    func = x + y;
    $display("x = %d, y= %d", x, y);
endfunction
```

Enabling Concurrent Export Task/Function Calls

Currently Simulation Acceleration allows only one execution occurrence of a specific export TF (task/function) at any given time, during run-time. If multiple calls to the same TF are made simultaneously, it results in an error, and stops Simulation Acceleration execution. Use the following options to enable multiple simultaneous occurrences of the same task, by specifying the name of the task/function and number of needed simultaneous occurrences of the respective task/function:

Use Model

V2Z Mode (UTF Option):

```
simxl_allow_concurrent -module <module-name> -tf {<task-function-list>}
[-count <number-of-occurrences>]
```

Where,

- **<module-name>:** Specifies the module or interface where the task/function resides.
- **<task-function-list>:** Specifies a space separated list of tasks/functions in the module/interface above.
- **<number-of-occurrences>:** Specifies the maximum number of simultaneous occurrences required for the specific task

Examples


```
simxl_allow_concurrent -module subdut -tf {dutTaskOne}  
simxl_allow_concurrent -module sub2 -tf {T1 F1} -count 2
```

V2VX Mode (VCS option):

```
-simxl=concurrent+<module name>+<tf name>[+count]
```

Examples (equivalent to above mentioned UTF commands)

```
-simxl=concurrent+subdut+dutTaskOne  
-simxl=concurrent+sub2+T1+2  
-simxl=concurrent+sub2+F1+2
```

Note:

- In both modes, count is optional, and defaults to a value of 3.
- In ZeBu mode, the list of TFs in one module, that have same concurrent count, can be specified in one command.
- If the TFs correspond to a different module, or have different concurrent counts, they have to be specified as a separate UTF command
- In VCS mode, each TF spec has to be specified as an individual suboption 'concurrent'. The values in the command are separated by a '+' and recognized by position.

Supported and Unsupported Behavioral Compiler Syntax

For information on the supported and unsupported Behavioral Compiler syntax, see the following subsections:

- [Supported Behavioral Compiler Syntax in Simulation Acceleration](#)
- [Unsupported Behavioral Compiler Syntax in Simulation Acceleration](#)

Supported Behavioral Compiler Syntax in Simulation Acceleration

The following behavioral compiler syntax is supported:

- [#delay Support](#)
- [Clocking Edge in for Loops](#)
- [Multiple Clocks or Edge Expressions in the Same Process](#)
- [Complex Edge Combination](#)
- [Combinations of Edge and Level Events](#)

- [Bounded Loops](#)
- [Unbounded Loops](#)
- [Wait Statements](#)
- [Named Events](#)
- [Behavioral Compiler Supports all SystemVerilog Data Types](#)
- [Clocking Blocks](#)
- [Fork/Join](#)
- [SystemVerilog Strings, \\$sformatf, %m are Supported](#)
- [Usage of Memory in Behavioral Code](#)

#delay Support

- Used to model clocks and resets with #delay.
- Use the `clock_delay` UTF command to specify the module #delay is used in. Note that the UTF command only applies to the module contents at one level and not to any of the instantiated modules.
- You can use this to model delays in behavioral code as well.
- You can use variable delays that can be reconfigured.
- There is no support for having `$readmem*` system calls embedded in initial statements that are used to implement clocking and reset behavior.
- Use the `clock_config -accuracy 16|24|32` UTF command to change bit-width to enable finer accuracy of clock delay precision.

To improve the performance degradation on the driver clock frequency, the `-perf_mode 2` option is introduced. This option replicates the RTL clock for SimXL transactors. The `-perf_mode 2` option is added to the `clock_delay -module` UTF command as follows:

```
clock_delay -module <module_containing_rtl_clock> -perf_mode 2
```

Example

```
module HW;
bit  aclk;
bit  aresetn;
initial begin
aresetn = 1'b0;
#5000;
aresetn = 1'b1; end
initial begin
aclk <= 1'b0;
```

```
forever begin
aclk <= 1'b1; #1000;
aclk <= 1'b0; #1000;
end
end
endmodule
```

Clocking Edge in for Loops

See the following example:

```
for (i=0;i<128;i=i+1) begin
@(posedge clk);
mem[i] = 0;
end
```

Multiple Clocks or Edge Expressions in the Same Process

See the following example:

```
always @(posedge clk) begin
a = 0;
@(negedge clk);
a = 1;
@(negedge reset);
a = 2;
end
```

Complex Edge Combination

See the following example:

```
always @(posedge clk or negedge clk2)
```

Combinations of Edge and Level Events

See the following example:

```
always @(clk or reset or posedge sig)
```

Bounded Loops

See the following example:

```
initial
for (int i=0; i< 5; i++) begin
c <= c + 1;
end
```

Bounded loops are unrolled if there are no synchronizations between HW and SW within the loop. For example, if there is a `$display` in the loop, it is not unrolled.

Unbounded Loops

See the following example:

```
initial
while (1) begin
@(posedge clk1)
c <= c + 1;
end
```

Wait Statements

To enable correct wait behavior, see the following example:

```
wait(clk);
```

Named Events

See the following example:

```
event my_event;
initial -> my_event;
always begin
@(my_event);
a <= b;
end
```

Behavioral Compiler Supports all SystemVerilog Data Types

Behavioral Compiler supports all SystemVerilog data types specified above.

Clocking Blocks

The following limitation applies:

Cannot access signals in interface's clocking block using XMRs from SW.

Example

```
clocking cb @(posedge aclk);
default input #`SETUP_TIME output #`HOLD_TIME;
input aresetn ;
input awaddr ;
endclocking
initial begin
@(cb);
A <= cb.awaddr;
end
```

Note:

Simulation Acceleration reports a Not Yet Implemented (NYI) error message when a clocking block is used for the interface port in the HW hierarchy.

Fork/Join

Supports the following variations:

- fork/join
- fork/join_none
- fork/join_any

The following limitations apply:

- Disable is not supported.
- Process class for thread introspection and control is not supported.
- fork/join_none requires significant FPGA area and should be used carefully.

To enable fork/join, add the following VCS compile option: `-Xzemi3=fork_join`

Example

```
fork
do_A();
do_B();
join
```

SystemVerilog Strings, \$sformatf, %m are Supported

The support of `$sformat` and `$sformatf` system task/function calls on HW need the communication of string variables across the HW/SW boundary.

- The system function `$sformatf` returns a string argument.
- The system task `$sformat` has an output port of type string.

Use Model:

The support of `$sformat` and `$sformatf` can be enabled with the combination of below two switches.

- `-simxl=<systasks or the corresponding entry in the UTF file.>`
- `-zebu_string=<MAX_STRING_LENGTH>:-The -zebu_string variable length is limited to a maximum of 128 characters by default. To increase the limit of maximum characters allowed in string variables, -zebu_string can be modified to -zebu_string=<desired_maximum_capacity>. For example, -zebu_string=256, strings of character lengths 256 are needed to run the design`

If `-zebu_string` is not provided, Simulation Acceleration reports an error stating that the `-zebu_string` should be enabled. Specifying a Module That Should Not Be Processed by Behavioral Compiler

Specifying a module not to be processed by Behavioral Compiler, which is instantiated in a transactor, is done using the attribute (*ZebuIgnoreZemi3Processing*) for that module.

Alternatively, you can use the following switch in the VCS command script:

```
-Xhwcosim=enable_zemi3_ignore vcs
```

This switch enables VCS to mark all modules under a transactor to skip zemi3 compiler if there is no task/function in the modules. This option yields better compilation results and lifts the limitation of using hardware top as a transactor in the DUT.

It is a recommended practice to do this for modules that contain code, which can be processed by the ZeBu synthesis tools.

In the example, both IIP_TOP and BOT2 are not processed by the Behavioral Compiler.

```
module xtor(input CLK1, input [31:0] dout1, input [31:0] dout2);
wire [31:0] din1;
wire [31:0] din2;
integer i ;
task waitn (input bit [31:0] cycle ) ;
begin
for (i=0 ; i < cycle ; i= i+1 ) begin
@(posedge CLK1) ;
end
end
endtask

IIP_TOP iip_top(dout1, dout2, din1, din2);
SUB1 s1(dout1, din1, din2);

endmodule

(* ZebuIgnoreZemi3Processing *)
module IIP_TOP(input [31:0] dout1, input [31:0] dout2, output bit [31:0]
din1, output bit [31:0] din2);
assign din1 = dout1 + 100;
BOT2 b2(dout2, din2);
endmodule

module BOT2(input [31:0] dout2, output bit [31:0] din2);
assign din2 = dout2 + 1000;
endmodule
```

Usage of Memory in Behavioral Code

By default, memory is considered as a state element. Behavioral compile maintains a copy of a state element in a variable suffixed with '_0' for every state variable. The variable samples the original state variable every driver clock cycle. This mechanism, causes the memory to be bit blasted and it is not mapped to a ZeBu memory.

Use the following UTF command:

```
memories -zmem -instance {<xtor name>.<memory name>}
```

Where:

- `<xtor-name>`: Specifies the name of the transactor module
- `<memory-name>`: Specifies the name of the memory instance within the transactor

Memory inferencing is required to keep the size of the generated logic for Transactors instantiated many times manageable. During VCS elaboration, messages are reported in the log file. These messages specify the memories inferred and memories not inferred. When memory is not inferred, it is important to review the messages in the VCS elaboration log file and correct the behavioral code accordingly, if memory inferencing is required.

Typically, the VCS elaboration log file messages are reported for signals which are above `memSizeThreshold`, but you might not be intending to infer a memory for them. If you want them to be a memory, the read-write access pattern of the signals should match a typical memory usage.

The following lists examples of messages reported in the VCS elaboration log file.

Message 1

```
Memory memA is skipped as could not extract write port for this  
sequential vector signal.
```

Reason for Message

`memA` is one-dimensional packed array and is written using full signal assignment. In this case, it is recommended to make this vector into a two-dimensional array with a reasonable bus size.

Memory inferencing is also difficult if there are multiple assignments to different bits in the indexed array. Collect all bits in a temporary variable and then assign the indexed array location with a single assignment of the temporary variable.

Message 2

```
Memory memB is skipped as this sequential signal has neither reset  
sequence nor write port
```

Reason for Message

This is a generic message for any type of memory candidate signal that has a write pattern from which neither write-ports nor reset condition can be successfully extracted. This occurs when it is not obvious how the condition for writing enables the location by the given address to get written.

Message 3

Memory memC is skipped as concat expression couldn't be separated to extract port writes.

Reason for Message

This message appears if the next state of the sequential signal has a concatenation-based assignment and write ports could not be extracted from it. This occurs if the implementation does not use an assignment to an array with an index but uses a full array assignment using concatenation.

Message 4

Memory memD is skipped as it is a dangling mem.
Source Node for dangling: _ASN_66

Reason for Message

This message appears when write-port or reset is extracted correctly but this memory is read fully in an assignment. In general, if a memory can be accessed without going through read ports, it is marked as dangling memory.

The following is an example that causes a dangling memory:

```
assign out = memD;
```

The following is the expected read pattern:

```
assign out1 = memD[raddr];
```

Message 5

Memory memE is skipped as it is a dangling mem.
Source Node for dangling: Child instance inport

Reason for Message

The reason for this message to appear is similar to that of [Message 4](#). In this case, the memory is completely passed to input port of a child instance. Therefore, the whole memory is read completely.

Unsupported Behavioral Compiler Syntax in Simulation Acceleration

The following SystemVerilog constructs are not supported by the Behavioral Compiler:

- SystemVerilog semaphores are not supported
- SystemVerilog mailboxes are not supported

- SystemVerilog classes are not supported. Note passing a class object as a formal argument is supported.
- SystemVerilog dynamic data types are not supported (for example, dynamic arrays, associative arrays, queues, and so on).
- Usage of UVM `config_db` to assign physical interfaces to virtual interfaces is not supported in HW.

3

Scheduling Semantics and Race Conditions

The Simulation Acceleration time-coupled mode guarantees that advancement of time on the SW side is consistent with advancement of time on the HW side. Event scheduling between HW and SW is essential to maintain consistency between a VCS built model run and a V2Z built model run. For more information on V2Z, see V2Z Mode.

Scheduling semantics of code running on SW side is compliant with the IEEE Std 1800-2012 LRM. Scheduling semantics of code running on HW side is based on a cycle-based semantics that is followed by tools like Emulation and Formal Verification driven by design clocks for non-transactor (for example, DUT) code and driven by a ZeBu driver clock (that is, `uclock`) for transactor.

Implementation of scheduling semantics for the handshaking between HW and SW has assumptions made to get better performance and to avoid other implementation limitations. In this clause, we try to specify what assumptions have been made to enable the user to ask the right questions when observing behavior that is not expected due to race conditions.

In the following clauses, race conditions are discussed that occur during:

- Initialization time
- SW/HW communication using signals
- SW/HW communication using subroutines

For more information, see the following topics:

- [Initialization Between SW and HW](#)
- [Resolving Clock/Data Race Conditions Between SW and HW](#)
- [Signal Handshake Between SW and HW](#)
- [Subroutine Handshake Between SW and HW](#)

Initialization Between SW and HW

Initialization is done in different stages. The considerations are as follows:

- Ordering between SW and HW initializations
- Initialization of static variables with their declaration time initial values
- Initialization of variables in initial block at time zero

You should look at key variables at end of time zero using debug facilities and rationalize the assigned values. Comparison with the results achieved with a simulation only run is essential.

Signal Handshake Between SW and HW

The SW side completes a delta cycle following the LRM Scheduling diagram as depicted in clause 4 of the LRM. The HW side then executes based on synthesis semantics a few driver clocks until it is ready to update the SW side. The SW code that is sensitive to intermediate events of HW signals are not observed.

Resolving Clock/Data Race Conditions Between SW and HW

When clocks are implemented in HW and SW uses them in event control statements, it is possible to indicate to the SW side that clock events coming from HW should be ordered to occur first from a SW perspective to avoid race conditions.

There are two orthogonal ways to specify the ordering.

1. Runtime option:

```
-simxl=clock_file:<filename>
```

Where, <filename> is the file that contains a list of hierarchical paths to clocks residing in HW. The file should contain one path per line.

2. Compile-time option:

```
-simxl=tbclock
```

This identifies clocks that appear as HW XMR when used in @ or wait statements in SW.

These options can be used both in the V2VX Mode and the V2Z Mode.

Subroutine Handshake Between SW and HW

For more information, see the following subsections:

- [Import Calls and Non-Blocking/Blocking Assignments](#)
- [Export Calls Scheduling](#)

Import Calls and Non-Blocking/Blocking Assignments

Calling a none time consuming subroutine from HW to SW assigns arguments to the subroutine like a blocking assignment. Then, control moves over to the SW side at the end of the delta cycle. If there is code executing after the subroutine call in the HW side, that code executes in a different delta cycle than it originally was expected to.

In the following example, there is a none time-consuming subroutine call (that is, observed), followed by a blocking assignment (that is, `data_out*`) from variables that have a non-blocking assignment (that is, `m_data_out*`) at the clock edge of the same delta cycle. In a regular VCS run, the variables that have been assigned in a non-blocking assignment have their previous values. When running with Simulation Acceleration they (that is, `data_out*`) retain their current values. The blocking assignments occur in a different delta cycle than the non-blocking assignments.

```
// HW interface
interface itf (input clk, rstn);
  reg data_out_vald;
  reg [0:7] data_out;
  always @(posedge clk) begin
    If (!rstn) begin
      data_out <= 0;
      data_out_vald <= 1'b0;
    else begin
      // NBA - Non-Blocking Assignment
      data_out <= data_out + 1;
      data_out_vald <= !data_out_vald;
    end
    task monitor();
    forever begin
      // Block assignment occurring in different delta cycle than NBA
      m_data_out_vald = data_out_vald;
      m_data_out = data_out;
      @(posedge clk);
      if (m_data_out_vald) begin
        // Calling none time consuming function in SW
        m_monitor.observed(m_data_out_vald, m_data_out);
      end
    end
  endtask
endtask
```

Export Calls Scheduling

When calling a function/task from SW to HW (that is, export), Simulation Acceleration does not immediately call function/tasks on the HW side because it would be very expensive. Instead, the SW wrapper initiates the actual HW call that is scheduled in preLER region (when control is given to HW). HW executes the task and control returns to SW. VCS schedules events in various runtime queues. preLER is pre Last Event Routine and is part of the Prepostponed region as described in the LRM. Runtime will trigger a call to ZeBu during this region.

4

Incremental Compile

Incremental compilation is implemented in the Simulation Acceleration flow. The Simulation Acceleration tool creates an incremental compile database to track all the HW and SW dependencies.

With incremental compilation, Simulation Acceleration compilation detects if the incremental compile interface database in the current compilation is identical or a subset of the database in a previous Simulation Acceleration compilation. If it requirement is met, the Simulation Acceleration compilation skips the ZeBu compilation.

Incremental compile enables separating SW compile (for example, UVM Testbench) from HW compile (for example, DUT and Transactors). Therefore, changes to SW do not cause a recompilation of the HW, which can take significant compile time. SW only compile uses existing VCS fast incremental compile capabilities.

The incremental compile database is created with no additional **zCui** options and it is tracked there on.

For more information, see the following subsections:

- [Modes for Performing Incremental Compile with zCui](#)
- [Support for \\$hw_force ECO](#)
- [Key Messages in Simulation Acceleration Incremental Compile](#)

Modes for Performing Incremental Compile with zCui

In Simulation Acceleration, there are two modes for performing incremental compile with **zCui**:

- **First Mode:** This is a regular compile. This is the current conservative mode and likely results in a HW recompile.
- **Second Mode:** In this mode, either of the following are possible:
 - **User-aware incremental compile:** Generates a `simv` based on the current HW database.
 - **Shared HW database:** Generates a reference DB (`zcui.work`) required for incremental compile. The shared HW database is a clean compile `zcui.work` database.

You need to use the `--simxlSharedDB` option to pass shared HW database during incremental compile. A new `zcui.work` database is created with the `zCui -w <DB_name>` option. The shared HW database and the `DB_name` should be different names. Otherwise, an error message is displayed.

To enable multiple different `simv`'s and `DB_name`'s based on a single `SharedDB` database, use the following command:

```
zCui --simxlSharedDB <reference DB> -w <DB_name> <additional zCui options>
```

If the `-w` option is not specified, the name of the generated database is `zcui.work` by default.

Incremental compile requires the SimXL SW/HW interface to be identical or a subset of the shared DB as determined by the SimXL compilation. If there is an error, the design requires a first mode compilation.

At runtime, pass the incremental compile database using the following command:

```
./simv +zebu.work=<DB_name>
```

It is recommended to retain the original `simv/simv.daidir/csrc` by creating a new directory and perform an incremental compile in the new directory.

Note:

SimXL supports the renaming of `simv` and `simv.daidi` using the VCS `-o` switch. The `csrc` is always overwritten.

Note:

There is no difference between how the two last modes are used, the difference is in how the users arrange their testbench environment. There are two kinds of use models:

- 1) Users that have a single `zcui.work` database that they share between different test cases, knowing that the test cases do not have any HW dependencies and they only require SW compiles.
- 2) Users that have a `zcui.work` for every test and they need the incremental compile feature for each of the tests separately.

Support for \$hw_force ECO

Simulation Acceleration saves you from a complete HW recompilation when there is no change in HW and there is no change in communication or if communication between previous and current compile has changed such that current is subset of previous. To further reduce the need to perform a complete HW recompilation, support is added for situations where new force statements (`$hw_force`) were added in the testbench to force HW signals.

For more information, see the following subsections:

- [Prerequisites](#)
- [Use Model](#)
- [Limitations](#)

Prerequisites

Before using this feature, address the following prerequisites:

- You must have a successful first compilation.
- In the first compilation, you must specify the following UTF command to enable ECO support for the second compilation:

```
eco -reserve_fwc 1024 -reserve_qiwc 0 -reserve_force 128
```

Use Model

The following use model describes the flow after you have added new XMRs to the testbench. These would include statements to force HW signals.

Perform the following steps:

1. Delete or rename the existing `eco.work`.
2. Run the second compilation with the following VCS options:

```
-simx1=no_dut_recompile,incr_comp_eco_force
```

Simulation Acceleration detects the newly added XMRs and provides guidance on using zECO support to compile the newly added XMRs.

1. To use zECO, run **zCui** with the `-d` option.

The `hwcosim_eco.tcl` file is saved in the current working directory and the ZeBu database is saved as `eco.work`. VCS automatically runs the zECO command.

1. Run `simv` with the following command:

```
simv +zebu.work=eco.work/zebu.work
```

Limitations

The following limitation exists:

- The `-reserve_force_assign` **zCui** option is not supported during the second compilation.

Key Messages in Simulation Acceleration Incremental Compile

The following sections describe the different scenario occurrences when using the two operation modes:

- [No Warnings Reported When Generating simv](#)
- [Warnings Reported While Generating simv](#)
- [Communication Changes Caused Warnings While Generating simv](#)
- [New simv Cannot be Generated](#)

No Warnings Reported When Generating simv

- Testbench source code might have changed.
- There is no change in DUT source code.
- There is no change in communication facilities.
- Testbench these changes are not reported.

Warnings Reported While Generating simv

- Testbench source code might have changed.
- DUT source code has changed.
- There is no change in communication facilities.
- VCS lists the DUT changes for you to review.
- The DUT change report indicates the modules that have been changed, added, or removed.
- The exact DUT constructs that have changed are not be reported.
- You can review the module change report to determine if the changes are important for the test.
- You can decide if this `simv` can work without HW recompile.
- Testbench changes are not reported.
- Review the warning messages related to changes in DUT.

Communication Changes Caused Warnings While Generating simv

- Similar to [Warnings Reported While Generating simv](#), but there are changes in communication.
- The exact communication changes are reported.
- The changes in communication are subset of previous communication and therefore Simulation Acceleration can recreate the necessary Verilog for the missing communication based on previous communication database.
- Review warning messages related to changes in DUT and communication.

New simv Cannot be Generated

- Communication has changed such that the changes are new or superset of previous communication.
- The exact communication changes are reported.
- `simv`, if generated, does not work without HW recompile.
- Error is reported and `simv` is not generated.

5

SVA Assertion and Coverage Support

In Simulation Acceleration, to enable functional coverage and assertions at compile time, use the following UTF commands:

```
coverage -enable true // Enables covergroups specified in HW
assertion_synthesis -enable ALL // Enables assertions and coverage
properties specified in HW
```

For more information, see the following topics:

- [Enabling Functional Coverage and Assertions at Runtime](#)
- [Controlling SVAs at Runtime](#)

Enabling Functional Coverage and Assertions at Runtime

To activate functional coverage and assertions at runtime, add the following options to the `simv` command:

```
-simx1=enable_dut_fcov,enable_dut_sva
```

Coverage data is generated in the `simv.zebu.vdb` file.

No special options are required to enable functional coverage and assertions in SW.

To merge coverage data from HW with coverage data from SW, use the following URG command:

```
urg -dir simv.vdb simv.zebu.vdb
```

Controlling SVAs at Runtime

The SVAs in the testbench side are enabled or disabled using existing **vcs** SVA controls, while the SVAs in the HW side are controlled by using runtime `hwcosim` options. You can perform the following activities:

- **Enabling SVAs During Runtime**

By default, the SVAs in the HW side are disabled at runtime in Simulation Acceleration, however you can enable it by using the following options:

```
-simx1=enable_dut_sva
```

Specific SVAs in the HW can be disabled by passing an Assert file along with the `disable_dut_sva` option. In the Assert file, state the names of the assertions that you want to disable.

```
simx1=disable_dut_sva,assertfile:<filename>
```

- **Enable Postprocessing of SVA**

You can enable postprocessing of the SVAs in the HW side using the following options:

```
-simx1=post_process_sva,dumpdir:<dir.zsva>
```

The results are outputted in the ZTDB file, which can be converted to text format using the `zSvaReport` utility.

- **Handling Fatal Assertions**

Specify the following option to stop emulation and simulation when a fatal assertion (`$fatal/$error`) failure is detected on the SVA in the HW side:

```
-simx1=stop_on_svafailure
```

On stopping, a UCLI prompt is available for further debugging.

- **Change Report Severity Level**

By default, only SVAs without an action block or with an action block containing a `$error` or `$fatal` display a message on the screen when there is a failure.

By default, messages are reported on the screen when there is failure for the following SVAs:

- SVAs without an action block
- SVAs with an action block containing a `$error` or `$fatal`

The minimum level of severity from which messages are displayed can be changed using:

```
-simx1=set_report_severity,level=<Severity_level>
```

Where, <Severity_level> can be one of the following:

- ZEBU_SVA_Failed_Fatal (highest severity)
- ZEBU_SVA_Failed_Error
- ZEBU_SVA_Failed_Warning
- ZEBU_SVA_Failed_Display (lowest severity)

Example

When the severity argument is set to `ZEBU_SVA_Failed_Warning`, assertions that have an action block with `$warning` and higher severity (`$error` and `$fatal`) are displayed. This applies only on SVAs in the HW side.

6

Generating the Profiler Report

As you begin to achieve runtime functionality with a Simulation Acceleration V2Z environment (V2Z Mode), it is important to start becoming performance and functional aware.

Functional awareness is handled by making sure test passes as expected and matches the golden data. Performance awareness is required to ensure that the runtime environment is on the right track for achieving the required performance goals.

Simulation Acceleration [Communication] Profiler is a tool to aid users in Simulation Acceleration performance analysis. It generates a Simulation Acceleration Profiler report at runtime that consists of the details of communication between HW and SW during simulation acceleration.

Simulation Acceleration Profiler report includes the number of communication events between HW and SW and the amount of data transferred due to these events. It also reports the time consumed in different types of events in the Simulation Acceleration run. Frequent and expensive communication between HW and SW can indicate performance bottlenecks.

By analyzing the profile report, you can leverage this information to make appropriate changes to the environment to achieve the intended performance goals. These changes might involve modifying code or moving the code blocks into the HW.

For more information, see the following topics:

- [Enabling Profiler](#)
- [Profiler Report Sections](#)
- [Example: Communication Profiler Report](#)

Enabling Profiler

To enable Profiler, use the Simulation Acceleration runtime options described in the following table.

Table 3 *Runtime Options to Enable Profiler*

Runtime option	Description
<code>-simxml=profile</code>	Enables Profiler. This command is required to write TS Interface information in the report.
<code>-simxml=profile,profile_log:myrep.txt</code>	Renames the report generated. By default, the report is generated in a file called <code>simxmlProfile.txt</code> .
<code>-simxml=profile,nheavy:<n></code>	Specifies the number of expensive steps. By default, 10 expensive steps are printed. Max value of <n> is 100.

The following UTF command is also required at compile-time to write TS Interface information in the report:

```
ztopbuild -advanced_command {set_simxml_opt -enable_profile yes}
```

Note:

Multiple options require only one specification of `-simxml=profile`. There are no compile-time options needed.

Profiler Report Sections

The report comprises the following sections:

- Header
 - Tool Version
 - Start/End time
 - Clock frequency and edge count
 - Computed emulation time
- Time consumption
 - Simulation Acceleration Elapsed time
 - Communication time (Data Preparation + Synchronization)
 - Testbench time
 - VCS wait time

- SW/HW Synchronization
 - Number of time steps requiring synchronization
 - Number of synchronization events
 - Number of time steps requiring time streaming
- Histogram, elapsed time distribution, and expensive time steps
- Data transferred and event counts
 - Each type of signal/memory
 - Each transactor
- Details of data transfer (for individual signals, memories, tasks and functions)

Data preparation time for most expensive times is broken down into following:

- PIO data packing
- TF data packing
- Signal Writeback time
- Signal Flushing time
- Memory Readback time
- Memory Writeback time
- Memory Flushing time

Details of the time are more clear in the translog report.

Example: Communication Profiler Report

In the following example report, the following applies:

- The numbers in tables above are representative and not necessarily related.
- By default, 10 expensive steps are printed. Use the following suboption to change the default setting:

```
nheavy:<n>" [max of 100]
```

Communication Profiler Report

```
SimXL Communication Profiler Report  
=====
```



```

+-----+
| VCS Version          | O-2019.06-Alpha_Full164 (ENG)
| VCS Build Date       | Build Date = Aug 16 2018 07:22:59
+-----+
| SimXL Start Time     | Wed Aug 29 11:37:23 2018
| SimXL End Time       | Wed Aug 29 11:37:23 2018
| Mode                 | Zebu
| VCS Simulation Time  | 200
+-----+
| ZEBU Version         | O-2019.06-Alpha
| Zebu work path       | [full path: zebu.work]
| ZRDB Id              | 55124962
| Driver Clock Frequency | 6250 kHz
| Driver Clock Cycles  | 1124948
| Timestamp Clock Cycles | 20
+-----+

```

1. Summary:

=====

1.1 Elapsed time

NOTE: %time is the percentage of Elapsed Time.

The breakup of Communication Time indicates % of Communication Time/% of Elapsed Time.

NOTE: HW time is not reported here, as the primary purpose of this report is the identification and analysis of communication and testbench bottlenecks.

```

+=====+=====+=====+
|                                     | Time Spent | %time |
+=====+=====+=====+
| SimXL Elapsed Time                 | 4.960 ms  | 100    |
+=====+=====+=====+
| Test Bench Time                    | 1.166 ms  | 23.5   |
+=====+=====+=====+
| Communication Time                 | 3.795 ms  | 76.5   |
+=====+=====+=====+
| CoEmulation Sync Time              | 0.000 ms  | 0/0    |
+-----+-----+-----+
| Synchronization                    | 3.765 ms  | 99.2/75.9 |
+-----+-----+-----+
| PIO data packing                    | 0.026 ms  | 0.7/0.5 |
+-----+-----+-----+
| TF data packing                     | 0.004 ms  | 0.1/0.1 |
+-----+-----+-----+
| Signal Writeback                    | 0.004 ms  | 0/0    |
| Signal Cache Flush                  | 0.004 ms  | 0/0    |
+-----+-----+-----+
| Memory Readback                     | 0.3 ms    | 0/0    |
| Memory Writeback                    | 1.4 ms    | 0/0    |
| Memory Cache Flush                  | 2.06 ms   | 0/0    |
+=====+=====+=====+

```

```

Zebu initialization time
=====+=====+
|ZEBU Load Time           |      4.139 s |
|=====+=====+
|ZEMI3 Init Time          |      0.130 s |
|=====+=====+
|ZEMI3 Close Time         |      0.005 s |
|=====+=====+

```

1.2 Synchronization events

```

-----
Number of simulation time steps requiring synchronization: 20
Number of SW/HW synchronization events: 53
Number of simulation time steps requiring time communication only: 0
Number of simulation time steps in SW time streaming mode: 0
Number of synchronization events in SW time streaming mode: 22
Average time spent per synchronization event: 0.013 ms
Average number of time steps per second: 4.98

```

```

Number of simulation time steps in SW time streaming mode: 0
Number of synchronization events in SW time streaming mode: 22

```

1.3 Histogram:

```

-----
| Sync Calls | Time Step count |
|-----|-----|
|          >8 |             14 |
|           6 |              8 |
|           3 |             21 |
|           1 |              1 |

```

1.4 Elapsed time distribution:

```

-----
| Time Range | Time step count |
|-----|-----|
| 10.000 ms - 100.000 ms |             2 |
| 1.000 ms - 10.000 ms |             2 |
| 0.100 ms - 1.000 ms |            30 |
| 0.010 ms - 0.100 ms |             7 |

```

1.5. Most expensive Time steps:

```

-----
[Sorted by time in V2Z, by no of calls (i/p + o/p + tf, sync) in V2VX]
| Timestamp | Synchronization | Input | Output |
| Task/Function | Elapsed |
| (ns) | calls | changes | changes |
| calls | Time | | | |
|---|---|---|---|---|
|-----|-----|-----|-----|-----|
|          110 |          3 |          1 |          2 |
|          1 | 0.023 ms |          1 |          2 |
|          70 |          3 |          1 |          2 |
|          1 | 0.014 ms |          1 |          2 |

```

Chapter 6: Generating the Profiler Report

Example: Communication Profiler Report

		190		3		1		1
	1	0.013 ms						
		150		3		1		1
	1	0.014 ms						
		30		3		1		1
	1	0.022 ms						
		0		3		1		2
	0	0.067 ms						
		130		3		1		1
	0	0.015 ms						
		100		3		1		0
	1	0.015 ms						
		90		3		1		1
	0	0.014 ms						
		40		3		1		0
	1	0.130 ms						

1.6 Transactor Information:

 Number of transactors: 9
 Number of zemi3 xtor threads: 4

2. Event summary by type:

=====

Signals:

Type	Total Changes	Total data
-----	-----	-----
Input	21	52 b
Output	14	448 b

InOut Signals:

Type	Total Changes	Total data
-----	-----	-----
Input	3	192 b
Output	11	352 b

XMRs:

Type	Total Changes	Total data
-----	-----	-----
SW->HW	23	736 b
HW->SW	6	192 b

Procedural Writes:

Type	Total Changes	Total data
-----	-----	-----
SW->HW	15	45 b

Forces:

Type	Total Events	Total data
-----	-----	-----
Force	7	224 b
Release	1	32 b

Chapter 6: Generating the Profiler Report

Example: Communication Profiler Report

Memory:

Type	Total Calls	Total data
Read	6	48 b
Write	4	24 b

Task/Function communication data:

[for user transactors, and XMR TF calls]

[Import is HW->SW call, Export is SW->HW call]

HW->SW	SW->HW	HW->SW	SW->HW	Instance Name
Calls	Calls	Total Data	Total Data	
5	2	128 b	160 b	tb.u_duv

3. Details of Data Transfer:

3.1 Details of signal changes:

Total Changes	Width	ID	Signal name
Inputs:			
20	1 b	0	tb.u_duv.clk
1	32 b	1	tb.u_duv.din
Outputs:			
11	32 b	3	tb.u_duv.counter
3	32 b	2	tb.u_duv.douts
InOut: SW->HW:			
3	64 b	2	tb.u_duv.counter
InOut: HW->SW:			
11	32 b	2	tb.u_duv.counter
XMR: SW->HW:			
23	32 b	7	tb.u_duv.u_sub.wval
XMR: HW->SW:			
3	32 b	5	tb.u_duv.u_sub.tb.wdin
3	32 b	6	tb.u_duv.u_sub.val
Procedural Write: SW->HW:			
5	3 b	1	tb.u_hwttop.u_duv.A_0
5	3 b	3	tb.u_hwttop.u_duv.A_1
5	3 b	5	tb.u_hwttop.u_duv.A_2
Forces:			
Force Calls	Release calls	Width	ID
7	1	32 b	4
			tb.u_duv.u_sub.val

3.2 Details of Task/Function communication:

NOTE: For internal communication:

HW->SW calls have 32 bits added to Input size [For HW->SW ID communication]

```

        SW->HW calls have 32 bits added to Output size [For HW->SW ID
communication]
        Class and Virtual interface method calls (indicated with '*')
have 64 bits added to Input size [For class/interface handle]
Instance 1: tb.u_duv.u_sub
| Total Calls|Direction| Input Size | Output Size | Task ID |
Task/Function name |
|-----|-----|-----|-----|-----|
|          5 | HW->SW |      32 b |         0 b |      1 |
tb.v_task
|          2 | SW->HW |      32 b |      32 b |      0 |
tb.u_duv.u_sub.setVin

Instance 2:
| Total Calls| Direction| Input Size | Output Size | Task ID |
Task/Function name |
|-----|-----|-----|-----|-----|
|-----|

```

3.3 Details of Memory data transfers:

```

| Reads | Writes | Read size | Write size | ID | Memory name
|-----|-----|-----|-----|---|-----|
|-----|
|      3 |      2 |      24 b |      12 b |  7 | top.dut.mem
|      3 |      2 |      24 b |      12 b |  8 | top.dut.mem1

```

4. Details of Data Transfer for SimXL infrastructure:

=====

NOTE: This section is intended for internal RnD usage.

4.1 Data Transfer Summary:

```

| Type | Total calls | Total data | Average calls |
Average data |
|-----|-----|-----|-----|-----|
| Input |      20 | 5.000 Kb |      1.000 |
256.000 b |
| Output |      11 | 2.750 Kb |      0.550 |
140.800 b |
| Force |       7 | 1.750 Kb |      0.350 |
89.600 b |
| SW->HW call|       2 |      128 b |      0.100 |
6.400 b |
| HW->SW call|       5 |      160 b |      0.250 |
8.000 b |
| Memory Read|       0 |         0 b |      0.000 |
0.000 b |

```

Chapter 6: Generating the Profiler Report

Example: Communication Profiler Report

```
|Memory Write|          0 |          0 b |          0.000 |
0.000 b |
```

4.2 Details of Signal data transfers:

XTOR ID	Task ID	Type	Size	Change count
1	0	Input	256 b	22
1	0	Output	256 b	11
2	0	Input	256 b	21
2	0	Output	256 b	3
3	0	Force	256 b	7

4.3 Details of Time communication:

Function Name	Type	Size	Call count
cosimTime_0_emuDelta	HW->SW	64 b	1
cosimTime_0_simDelta	SW->HW	64 b	1

4.4 Details of Control communication:

Function Name	Type	Size	Call count
cosimCtrlIOSync	HW->SW	32 b	21
cosimCtrlISync	SW->HW	32 b	32

4.5 Details of Task/Function Mask communication:

XTOR Name	ID	Type	Size	Call count
cosimPIO_1_	1	POMask	32 b	10
cosimPIO_1_	1	PIMask	32 b	20
cosimTF_2_	2	POMask	32 b	7
cosimTF_2_	2	PIMask	32 b	7

5. TS interface information:

```
=====
tb.CLKCTRL.rtl_clkrdy: VC#: 6251458, !Ready Cycle#: 85704063
tb.cosimConf_2_.zemi3_readyForClock: VC#: 3, !Ready Cycle#: 2
tb.cosimCtrl.zemi3_readyForClock: VC#: 1993107, !Ready Cycle#: 996553
tb.cosimFR_1_.zemi3_readyForClock: VC#: 2208021, !Ready Cycle#:
1196924
tb.cosimTime_0_.zemi3_readyForClock: VC#: 1986365, !Ready Cycle#:
993182
tb.u_duv.clk0: VC#: 1553044
tb.u_duv.clk1: VC#: 1911
tb.u_duv.clk10: VC#: 161441027
```

6. Additional Information:

```
=====
Number of simulation time steps in HW time streaming mode: 1
Number of synchronization events in HW time streaming mode: 76
===== End SimXL Communication Profiler Report =====
```

7

Debugging and Reports

For more information, see the following topics:

- [Debug Options for Communication Mechanisms](#)
- [Functional Debug in Simulation Acceleration](#)
- [Using simprofile with Simulation Acceleration](#)

Debug Options for Communication Mechanisms

By default, the logs show the packed value of all arguments in each task/function call. This is useful for run-to-run comparisons. If you want to see detailed information on the task function arguments, use the following compile-time option:

```
-simx1=enable_log_tf
```

To enable logging the communication transactions that occur between SW and HW, use the following runtime options.

Table 4 *Runtime Options for Logging Communication Transactions Between SW and HW*

Options	Description
<code>-simx1=translog,translog_start:10,translog_end:100</code>	Enable logging from 10 to 100 time units
<code>-simx1=translog,translog_file:hwlog.txt</code>	Enable logging and print messages in file
<code>-simx1=translog,translog_level:1</code>	Enable logging and set verbosity level

It is recommended to use these options before going to debug through waveform signals. This provides a high-level view of the interactions going on between HW and SW.

To analyze a specific timestamp to get a hint as to what is happening on the SW side, perform the following:

1. Create a shell script with run permission, "mon_stack":

```
#!/bin/bash

if [ -z "$1" ]; then
    echo "Usage: $0 pid"
    exit
fi
pid=$1

echo "Starting spying on $pid ..."
while [ -e /proc/$pid/status ]; do
    echo
    /bin/date --rfc-3339=ns
    echo "======"
    pstack $pid
    usleep 100000
done
```

2. Add the UCLI command to the VCS UCLI script or run from the UCLI prompt:

```
ucli% exec mon_stack [pid] > pstack_dump.log &
```

Run to the timestamp (Using the `run -absolute <timestamp>` UCLI command) before the slow timestamp. Execute this command and stop `simv` after the slow timestamp. The script "mon_stack" will stop after `simv` stopped.

3. Check `pstack_dump.log` to find repeated called API or any other hint from call stacks. Send, capture, or share with R&D for analysis.

Functional Debug in Simulation Acceleration

Simulation Acceleration debug attempts to mimic VCS UCLI-based debug methodology. The UCLI provides basic functionality:

- Outputting waveforms using Dynamic Probe/FWC/QIWC
- Ability to set breakpoints in testbench code.
- Ability to force or deposit signal in HW.
- Ability to set breakpoint on signal changes in HW.

This section describes the following Simulation Acceleration debug-related information:

- [Enabling UCLI Signal Access for ZeBu Compile](#)
- [Enabling UCLI Waveform Output for ZeBu Compile](#)

- [Supported UCLI commands in Simulation Acceleration](#)
- [Verdi Interactive Debug for Simulation Acceleration](#)
- [Waveform Outputting Using UCLI](#)

Enabling UCLI Signal Access for ZeBu Compile

UTF commands used to enable UCLI control of signals in DUT

- **Force/Release**

```
zforce -rtlname <Zebu_mapped_signal>
```

This enables force or release from the command line.

- **Deposit Signal**

```
zinject -rtlname < Zebu_mapped_signal>
```

This UTF command makes transformations to enable deposit signal values.

- **Read signal values**

```
probe_signals -type dynamic -rtlname < Zebu_mapped_signal>
```

This UTF command allows reading signal values in UTF commands.

This applies to dynamic signals required when read in UCLI, such as usage of the UCLI get command.

You can use SystemVerilog commands to enable UCLI control of signals in DUT:

- Usage of `$hw_read` in SystemVerilog code is required in the SW side to specify the signal used in the UCLI. Note that these signals cross between HW and SW on every update. This is required for stop UCLI commands.
- Usage of `$hw_force` in SystemVerilog code is required in the SW side to force/deposit the HW signal in UCLI.

Enabling UCLI Waveform Output for ZeBu Compile

- Waveform outputting must be specified in HW Verilog module to be able to output waveforms using FWC/QWIC.
- It is recommended to create all the FWC/QWIC groups in a separate module.
- Make sure the following commands are specified in the UTF file:

```
debug -waveform_reconstruction true  
debug -verdi_db true
```

- Use an initial block with a label to identify a group of signals with a name. This name is used in the `dump` commands as the "value set". The name of the signal group is the same name as the "value set".

```
module dumpvars ();

initial begin: fwc_waves_grp1
(*fwc*) $dumpvars(0, hw_top.inst.a);
end
initial begin: qiwc_waves_grp1
(*qiwc*) $dumpvars(0, hw_top.inst.b);
end
initial begin: fwc_waves_grp2
(*fwc*) $dumpvars(0, hw_top.inst.c);
end
initial begin: qiwc_waves_grp2
(*qiwc*) $dumpvars(0, hw_top.inst.d);
end
endmodule
```

Supported UCLI commands in Simulation Acceleration

See the following subsections:

- [get UCLI Command](#)
- [memory UCLI Command](#)
- [stop UCLI Command](#)
- [run UCLI Command](#)
- [force, release, show UCLI Commands](#)
- [restart UCLI Command](#)
- [System Information Commands](#)
- [dump Commands](#)
- [sniffer Command](#)

get UCLI Command

Obtain the value of a signal or variable

Syntax

```
get <signal> [-radix string]
```

Get the value of a signal *<signal>* in the format specified by the `-radix` option.

Example

```
get u_duv.ctrl -radix b
get u_duv.ctrl -radix d
get u_duv.ctrl -radix h
get u_duv.mem0.m -radix d # also works on memory
get u_duv.mem0.m[100\] -radix d
```

Limitation

To display signals on the HW side, you must specify them at compilation time.

Use Model

To display signals on the HW side, specify them at compilation time in one of following ways:

1. Use `$hw_read` system task call.
2. Use `UTF probe_signals` command.
3. Use `$display` or any other read access in SW.
4. XMR procedural access from the SW side.

memory UCLI Command

Load or write memory values from or to files, or initialize memory with a specified value.

Syntax

```
memory -read|-write -file <fname> [-radix <radix>] [-start start_address]
      [-end end_address]
```

Example

```
memory -read u_duv.mem0.m -file in_mem_bin.txt -radix b
memory -read u_duv.mem0.m -file in_mem_hex.txt -radix b
memory -read u_duv.mem0.m -file out_mem_hex.txt -radix hex
memory -write u_duv.mem0.m -file hello.txt
```

The following are the equivalents of the Verilog system task `$readmem`, `$writemem`.

```
memory -read -radix b -> $readmemb
memory -write -radix h -> $writememh
```

Use Model

To be able to read or write HW memory, enable access to the memory.

1. Use `$hw_read` and/or `$hw_write` system task call.
2. XMR usage in TB in procedural context.

stop UCLI Command

For `stop` command usage when Dynamic Trigger and Runtime Trigger, see the following subsections:

- [Simulation Acceleration Dynamic Trigger](#)
- [Simulation Acceleration Runtime Trigger](#)

Simulation Acceleration Dynamic Trigger

The following table shows dynamic trigger options that pertain to the `stop` command.

Table 5 *Dynamic Trigger Options for the stop Command*

stop Command	Description
<code>stop -expression <expr> -name <name></code>	Enables the specified <name>, where- <name> is the instantiation path- <expr> is the equation
<code>stop -command <proc> <name></code>	Enables execution of a <proc> Tcl procedure when the equation is true. Execute this command after the following command: <code>stop -expression <expr> -name <name></code>
<code>stop -enable <name></code>	Enables the specified stop condition <name>
<code>stop -disable <name></code>	Disables the specified stop condition. Here, <name> should be a notifier name or a trigger path.
<code>stop -is_dynamic <name></code>	Returns 1 if <name> refers to a dynamic stop, else 0 is returned.

Simulation Acceleration Runtime Trigger

Before using runtime triggers in Simulation Acceleration, make sure you have the following files:

- The `.cel` file that describes the FSM. The description is in the Complex Event Language (CEL).
- The `Dumpvars` file (FWC, QiWC) contains all the signals defined in the `.cel` file.

The runtime trigger options that pertain to the `stop` command are as follows:

- **List all configured stop conditions, including runtime trigger objects:** Specify the `stop` command with no options.

For example:

```
$ stop
1: -cel a -fid ZTDB0 -once -action d
2: -cel a -fid ZTDB0 -repeat -action d
```

- **Enable the runtime trigger:** Specify the following command:

```
stop -cel [-fid <fid>] [-action <TclProcName>] [-once] [-repeat <N> |
repeat_forever]
```

Where,

- `-action <TclProcName>`: Specify to execute a Tcl procedure when the stop condition is reached. The Tcl procedure must have the following format:

```
proc call {module RTLTimeStamp lastNotify} {
    [proc body]
}
```

Where,

- `module`: Represents the module name where the notification occurs.
 - `RTLTimeStamp`: Represents the RTL timestamp when the notification occurs.
 - `lastNotify`: If software notifier was enabled in repeat, the value for this parameter is 1 only for the callback for the last notification, otherwise the value is 0.
 - `-once`: Specify to run the CEL FSM once. This is the default behavior if `-repeat` or `-repeat_forever` is not specified.
 - `-repeat <N>`: Specify to run the CEL FSM for a specific number of times `<N>`.
 - `-repeat_forever`: Specify to continuously run the CEL FSM until you disable the runtime trigger.
- **Disable the runtime trigger:** Specify the following command:

```
stop -disable <moduleName>
```

To enable the runtime trigger, specify the following command:

```
stop -disable <moduleName>
```

Example: Using the Runtime Trigger in Simulation Acceleration

For FWC, you need to specify the `-add_value_set` option, as shown in the following snippet:

```
dump -file fwc.ztdb -type fwc
dump -add_value_set ZFWC_DFLT_GRP -fid ZTDB0
dump -disable -fid ZTDB0
stop -cel ../run_zebu/pattern.cel -fid ZTDB0 -once
dump -enable -fid ZTDB0
run 500ns
```

For QiWC, you do not need to specify the `-add_value_set` option.

An example of the callback procedure is as follows:

```
proc callback {module rtl_time is_last_notify} {
    echo "***** CALLBACK *****"
    echo "MODULE           : $module"
    echo "USER CYCLE         : $rtl_time"
    echo "LAST NOTIFY        : $is_last_notify"
    echo "*****\n"
}

dump -file fwc.ztdb -type fwc
dump -add_value_set ZFWC_DFLT_GRP -fid ZTDB0
dump -disable -fid ZTDB0
stop -cel ../run_zebu/pattern.cel -action callback -fid ZTDB0 -once
dump -enable -fid ZTDB0
run 500ns
```

To view the module name of the loaded CEL file, specify the following command:

```
stop -cel
```

run UCLI Command

To start a simulation run that continues up to a specified time or event.

Example

```
run -absolute 10
run -posedge clk
run -absolute 5ns
```

force, release, show UCLI Commands

Force or deposit a value on a signal or variable.

Use the `force -list` command to show UCLI forced signals.

Syntax

```
force <signal> <value> -deposit|-freeze  
force -list
```

If the `-deposit` option is specified, value is deposited on the current cycle to the signal `<signal>`.

If the `-freeze` option is specified, the value is frozen until released.

The `release <signal>` command releases a previously frozen signal.

Example

```
run 50ns  
force counter 32'd1 run 50ns  
release
```

To see if a signal is freezable or depositable, use the following command:

```
show -freezable|-depositable <signal>
```

This command returns if the signal passed as parameter is forceable, injectable, or writable.

Example

```
ucli% force dut_top.master_bfm_inst.wdata 0  
ucli% run 100  
100100 ps  
ucli% force -list  
dut_top.master_bfm_inst.wdata  
ucli% release dut_top.master_bfm_inst.wdata  
ucli% force -list  
No signal has been forced via UCLI.
```

Note:

- The signals that are forced with the `-deposit` option are not listed
- The signals that are already released by the `release` command are not listed
- It only lists the UCLK forced signals. The signals forced by Verilog system task are not listed

restart UCLI Command

To restart the simulation at time zero, use the following UCLI command:

```
restart
```

System Information Commands

The UCLI commands used to retrieve system information is as follows:

Table 6 *UCLI Commands to Retrieve System Information*

UCLI command	Description
<code>senv</code>	Display one or all <code>synopsys::env</code> array elements.
<code>senv value_sets</code>	Returns all value sets existing in design compilation.
<code>senv driver_clk_frequency</code>	Returns driver clock frequency in kHz.
<code>senv zebu_work</code>	Returns <code>zebu.work</code> path.
<code>senv time</code>	Returns current emulation time.

dump Commands

See the following `dump` commands:

- `dump -file <FILE> -type fwc|dynamic_probe`

Creates a waveform database directory named `<FILE>` for both SW and HW. Returns a file ID (that is, FID) that can be captured in Tcl. The subsequent commands use the FID. The default FID is "ZTDB0" and can be used directly without capturing it in a Tcl variable. In the waveform database directory, an FSDB file named `simx1.fsdb` is used to capture all SW signals.

- `dump -add <list of hierarchical paths> -depth <positive number> -fid <FID>`

Adds signals that are outputted to the provided FID. Specify the instance and depth from which the signals are to be taken. Depth 0 provides all signals in the provided instance and below.

- `dump -add_value_set <value_set> -fid <FID>`

Adds `<value_set>` to a given `<FID>`. Multiple value sets can be added as separate values but not as a list. Consider the following examples.

Correct Usage

```
dump -fid $fid -add_value_set {ZEBU_WAVES} {ZEBU_WAVES_DUMPPOINTS}
```

Incorrect Usage

```
dump -fid $fid -add_value_set {ZEBU_WAVES ZEBU_WAVES_DUMPPOINTS}
```


Value sets are specified using `$dumpvars` commands together with block labels.

- `dump -enable/-disable -fid <FID>`

Enables and disables outputting before starting or continuing simulation.

- `dump -close -fid <FID>`

Closes the file associated with the specified FID.

- `dump -flush -fid <FID>`

Flushes the contents into the file associated to the specified FID.

- `dump -load_selection [<path>] -fid <fid>`

Used to specify a selection database. If no path is provided, the following path is used:

```
./zcui.work/zebu.work/zrdb/csa_supports.zrdb
```

The `csa_supports.zrdb` database file is created by either of the following commands:

```
debug -waveform_reconstruction true
```

or

```
debug -all true
```

sniffer Command

Simulation Acceleration leverages the ZeBu Server debug technology called **Stimuli Capture and Replay** to rerun the design deterministically and to capture waveforms for a part (frame) of the run.

The Sniffer technology records the ZeBu state in frames and saves the stimuli. Each frame is a window of stimuli with the ZeBu State at the beginning of the stimuli window. To use the Sniffer technology, specify the `sniffer` UCLI command.

For more information on debug technologies, such as Stimuli Capture and Replay, `sniffer`, and `replay`, see the *ZeBu Server Debug Guide* and the *ZeBu Server Debug Methodology Guide*.

For more information, see the following subsections:

- [Compilation Setup](#)
- [Supported sniffer Options for Simulation Acceleration](#)
- [Setting the Runtime Generated Data Directory](#)
- [Use Model](#)

- [Limitations](#)
- [sniffer Examples](#)

Compilation Setup

The **Stimuli Capture and Replay** feature is enabled by default. To disable this feature, add the following UTF command in your UTF file: `debug -offline_debug false`.

For more information, refer to the **Compilation Setup for Stimuli Replay** section in *ZeBu Sever Debug Guide*.

Supported sniffer Options for Simulation Acceleration

The following options are supported for the `sniffer` UCLI command:

- `sniffer -start_frame`
- `sniffer -at_speed`
- `sniffer -config save_state_at_stop`
- `sniffer -config keep_frames <n>`
- `sniffer -config output_record`
- `sniffer -config no_memory_copy`
- `sniffer -auto_create <minutes>|<seconds>s|<minutes>m|<hours>h [-prefix <prefix>]`
- `sniffer -list [-dbs | <db_name>]`
- `sniffer -info [<name>]`
- `sniffer -status`

Setting the Runtime Generated Data Directory

By default, runtime generated data is stored in the directory called `emulation_data_<DATE>_<TIME>`, which is in the current working directory. To set a different location for the runtime generated data directory, use the following command:

```
simv -ucli -simx1=db_path:<path>
```

Where, `<path>` is the new directory name in the current working directory. For example, in the following command, `db2` is the new directory name:

```
simv -ucli -simx1=db_path:db2
```

Use Model

1. Create frame using the create option of the sniffer UCLI command:

```
sniffer -create <label_name>
```

2. Stop capture using the stop option of the sniffer UCLI command:

```
sniffer -stop
```

3. You can use **zRci** to replay the captured frames. In **zRci**, use the `replay` command to achieve this.

Example **zRci** script for `replay`

```
config zebu_work zcui.work_default/zebu.work/ config db_path db
replay -config enable_state_checks on
sniffer -import sniffer_good_backup # lists captured frame names
sniffer -restore 20190510_164812.F_1
config default_clock replay 75ns
exit
```

To view all options of the command, specify the following command:

```
replay -help
```

Limitations

To prevent abnormal behavior during replay, you need to explicitly call the `sniffer -stop` command before the simulation has completed.

sniffer Examples

In the following examples, the final label names are `sn_0`, `sn_10`, `sn_20`, and so on.

```
# Example 1
set idx -1
sniffer -create sn_[incr idx]
stop -relative 10ns -continue -repeat -command {
    sniffer -create sn_[incr idx]
}
# Example 2
sniffer -create sn_$now
stop -relative 10ns -continue -repeat -command {
    sniffer -create sn_$now
}
```

powermgt UCLI Command

Prerequisites

The `powermgt` command only works on an UPF design compiled with ZeBu Power Aware flow. To compile an UPF design, specify the UPF file in the VCS command line as follows:

```
vcs -upf <filename.upf> <vcs_options> <design_files>
```

For more details about the supported version of UPF file, see the "Power Management Script (UPF)" in *ZeBu® Power Aware Verification User Guide*.

When an UPF design is emulated in ZeBu PowerAware emulation flow using SimXL, a power domain switches on or off and pseudo random values are injected into its internal state elements such as registers, latches, and memories. To prevent data scrambling in simXL, simXL supports the following UCLI command so that data scrambling can be disabled:

```
powermgt -scramble -enable|-disable [-all|-registers|-memories]
```

where,

- `-scramble`: Enables or disables the scramble feature. Defaults to both registers and memory on enable/disable operations.
- `-all`: Selects both registers and memory for enabling/disabling the scramble feature.
- `-registers`: Selects registers for enabling/disabling the scramble feature.
- `-memories`: Selects memory for enabling/disabling the scramble feature.

Verdi Interactive Debug for Simulation Acceleration

Verdi interactive debug supports the Simulation Acceleration flow. With Simulation Acceleration interactive debug, you can debug both software (testbench) and hardware (DUT) in the Verdi interactive debug environment. Simulation Acceleration interactive debug also allows you to output testbench signals and DUT signals into one ZTDB file.

For more information, see the *ZeBu Verdi Integration Guide*.

Waveform Outputting Using UCLI

- Simulation Acceleration UCLI uses `dump` commands described in the previous section to control the outputting.
- The key differences between VCS and ZeBu output is as follows:
 - VCS does not have any special groupings to output waveforms.
 - VCS only requires read access enabled for outputting.

All the hierarchies and signals are available for outputting.

- ZeBu has the following separate modes for outputting:
 - **Dynamic Probe**: This mode can be used to output any hierarchy without compile directive, but is very slow.
 - **FWC**: This mode is fastest but requires significant hardware resources.
 - **QiWC**: This requires fewer HW resources and has intermediate performance.
- ZeBu outputs waveforms in a compact format called ZTDB:
 - ZTDB format must be converted into the ZWD format to be enabled in Verdi.
 - ZTDB only outputs register values and "combinatorial signals" are calculated on-the-fly by Verdi.
 - Unlike pure VCS, the waveform outputting should be properly planned before compiling the design.

For more information, see the following subsections:

- [Dynamic-Probe](#)
- [FWC Probe](#)
- [QiWC Probe](#)
- [Multi-ZTDB Output, Multi-ValueSet Capture](#)
- [Using Verdi for Debug](#)
- [Output ZeBu XTOR Internal Signals](#)
- [Waveform Expansion in Simulation Acceleration](#)
- [Waveform Outputting config slice for zSimzilla](#)

Dynamic-Probe

Advantages

- No compile directives required
- Full scope of design is available for dumping

Disadvantages

- Very slow outputting. Probably suitable for only few cycles for a large design.
- Recommend to run within a very short window using absolute run commands (see below)

- UCLI Tcl outputting example:

```
run -absolute 34000ns
set v_fid [dump -file rwc_wave.ztdb -type dynamic_probe -driverClk];
dump -load_selection -fid $v_fid;
dump -enable -fid
$v_fid;
run -absolute 34010;
dump -fid $v_fid -flush;
dump -fid $v_fid -close;
```

FWC Probe

Advantages

- This is fastest way to output probes.

Disadvantages

- Must be enabled at compile time
- Uses extensive HW capacity and can only be enabled for selected "essential signals"

Example

```
set v_fid [dump -file fwc_wave.ztdb -type fwc]
dump -add_value_set fwc_waves_grp1 -fid $v_fid
dump -add_value_set fwc_waves_grp2 -fid $v_fid
dump -enable -fid $v_fid
run 10us
dump -fid $v_fid -flush
dump -fid $v_fid -close
```

QiWC Probe

In V2VX mode, you can do the same waveform outputting and debug as in regular VCS flow.

In V2Z mode, you can output the waveform in **zRci** GUI. For details, see the *ZeBu User Guide*. Simulation Acceleration V2Z mode also allows you to use ZeBu QiWC waveform capture through UCLI. The steps are as follows:

1. In the DUT, add following source code to enable ZeBu QiWC waveform output:

```
initial begin : GROUP
(* qiwc *) $dumpvars(0, tb.dut);
end
```

2. Create a UCLI script to output the waveform.

```
//mydump_qiwc.tcl:
dump -file myqiwc.ztdb -type fwc
dump -add_value_set GROUP -fid ZTDB0
```

```
run 372685000
dump -close
```

3. Run the simulation with this UCLI script as shown in the following command:

```
% simv -ucli -i mydump_qiwc.tcl <other runtime options>
```

4. After simulation you can open and debug the outputted waveform with Verdi by using the following command:

```
% verdi -emulation --input <your dumped ztdb directory> --zebu.work
<your zebu.work directory>
```

Outputting Signals of HW Module Instances That Were Moved to SW

By default, module instances that were moved from HW to SW are not outputted. Outputting request needs to be specified explicitly using the following UCLI command:

```
dump -add {<module-instance> } -fid <File-Identifier-Name>
```

Example

```
dump -add {top top.dut.subdut.tb_inst1} -fid ZTDB0
```

Note:

These signals are sampled into the `simx1.fsdb` file residing in the waveform database directory associated with the FID, named ZTDB0.

Multi-ZTDB Output, Multi-ValueSet Capture

Simulation Acceleration supports multi-ZTDB output through UCLI. Several FWC and QiWC waveform files can be captured simultaneously as long as they do not share any Value-Set.

Simulation Acceleration also supports capturing multiple FWC and QiWC into the same output. For readback capture, only one dynamic-probe waveform can be captured at a time.

Example

```
set fid [dump -file dump.ztdb -type fwc]
dump -add_value_set FWC -fid $fid
dump -add_value_set QiWC -fid $fid
dump -enable -fid $fid
run 100ns
dump -close -fid $fid
quit
```

Output FWC/QiWC Waveform Only on Time Advance

You can use the `sample_on_time_adv` option to reduce the time stamp or raw data file size in the ZTDB file as pound zero (#0) occurs in the design. This option enables sampling on time-advance while FWC or QiWC is outputted. Converting the reduced files saves waveform conversion time to increase convert performance.

Prerequisite

The Verilog RTL file must contain #0.

Procedure to Enable

1. Run `simv` with the `sample_on_time_adv` option, as shown in the following:

```
simv -ucli -simx1=sample_on_time_adv
```

2. Output FWC/QiWC using UCLI. See the [Waveform Outputting Using UCLI](#) section.

If the sample-on-time-advance run is successful, a `config_timestamp` that contains sample information in the ZTDB directory is created, as shown in the following:

```
SAMPLING_ONLY_ON_TIME_ADVANCE TRUE
```

Limitations

- Supported only for ZeBu Server 4
- The `-simx1=disable_alt_exec_flow` VCS command-line options is not supported

Using Verdi for Debug

- [Verdi Interactive Debug for Simulation Acceleration](#) is supported. See the *ZeBu Verdi Integration Guide* for more information.
- Use "`-kdb -lca`" VCS option to output Verdi KDB database like VCS.
- Open the Verdi database with the ZTDB file using the following command:

```
verdi -nologo -emulation --zebu.work <ZEBU_WORK> --input <ZTDB_FILE>  
<other verdi command line options>
```

This opens Verdi to be run with the iCSA engine.

Output ZeBu XTOR Internal Signals

When using the subroutine-based Simulation Acceleration flow, you might want to output XTOR internal signals. This is useful, for example, when you want to check XTOR key FSM signals "`imp_state_net_*`" to debug an XTOR behavior. You can output these

signals through the Simulation Acceleration `dynamic_probe` by performing the following steps:

1. Run `zSelectProbes` to select the signals to output.



1. Select **Save selection** from the **File** menu to save the selects.
2. Exit `zSelectProbes`.
3. Specify the following UCLI commands to output the signals:

```
set fid_dyn [dump -file mydyn.ztdb -type dynamic_probe -driverClk]
dump -load_selection -fid $fid_dyn
dump -enable -fid $fid_dyn
run 8135000
dump -flush -fid $fid_dyn
dump -close $fid_dyn
```

After running `zSelectProbes`, besides outputting the selected signals through UCLI command line, you can output the selected signals in `zRci`.

Waveform Expansion in Simulation Acceleration

After the ZTDB file is outputted, you can use the ZeBu Waveform Expansion utility, `zSimzilla`, to get the ZWD file for debugging with Verdi. For more information, see the *ZeBu Debug Guide*.

Example

The following command writes FSDB files from `myqiwc.ztdb` to the `dut_fsdb` directory. You can then debug the design with the generated FSDB files.

```
%zSimzilla --work zebu.work --ztdb myqiwc.ztdb dut_fsdb --timescale 1ns  
%verdi -ssf dut_fsdb.vf
```

For XTOR internal signals outputted through dynamic-probe, you can only convert ZTDB to FSDB. There is no waveform expansion available for these internal signals. The `zSimzilla` command is as follows:

```
%zSimzilla --capture-only --work zebu.work ...
```

Waveform Outputting config slice for zSimzilla

Compile-Time Prerequisite

To use **zSimzilla**, make sure the following command is specified in the UTF file:

```
debug -waveform_ reconstruction true
```

Outputting ZTDB Type QiWC or FWC

Specify the following command for outputting ZTDB of type QiWC and FWC:

```
config waveform_capture_slicing  
(auto|<seconds>[s]|<size>MB|<samples>total_samples|<slices>slices|{<samp  
les>total_samples,<slices>slices})
```

Outputting ZTDB Type dynamic_probe

Specify the following command for outputting ZTDB of type `dynamic_probe`:

```
config dynamic_probe_slicing  
(auto|<seconds>[s]|<samples>total_samples|<slices>slices|{<samples>total  
_samples,<slices>slices})
```

Example 1

This example shows how to enable `config slice` feature.

Specify the following to enter interactive mode:

```
../simv -ucli
```

In readback, specify the following:

```
config dynamic_probe_slicing auto  
dump -file rb.ztdb -type dynamic_probe  
dump -load_selection -fid ZTDB0  
run
```

For FWC, specify the following:

```
config waveform_capture_slicing auto  
dump -file fwc.ztdb -type fwc  
dump -add_value_set {VAL} -fid ZTDB0  
run
```

Run zSimzilla to convert ZTDB to ZWD by entering the following commands:

```
zSimzilla --zebu-work ../zcui.work/zebu.work \  
  --ztdb fwc.ztdb \  
  
  --zwd hello.zwd  
nWave hello.fsdv.vf
```

Example 2

This example shows some syntactical use of `waveform_capture_slicing` and `dynamic_probe_slicing`:

```
config waveform_capture_slicing 3MB  
config waveform_capture_slicing 3s  
config dynamic_probe_slicing 3s  
config dynamic_probe_slicing 100total_samples  
config dynamic_probe_slicing 10slices  
config dynamic_probe_slicing 100total_samples,10slices
```

Using simprofile with Simulation Acceleration

The Unified Simulation Profiler (simprofile) is supported in Simulation Acceleration.

For more information, see the following subsections:

- [Viewing Timeline Using simprofile](#)
- [Viewing Delay Events Using simprofile](#)

Viewing Timeline Using simprofile

With simprofile, it is possible to get a breakdown of the testbench code and see how much time is used by the Simulation Acceleration communication overhead. To enable, use the regular simprofile compile time and runtime options.

It is also possible to get a profile over time to see how the testbench is a bottleneck at different timepoints throughout the run. This feature is the simprofile timeline.

To enable, perform the following:

1. Specify the following compile option:

```
-simprofile
```

2. Specify the following runtime option:

```
-simprofile timeline
```

Viewing Delay Events Using simprofile

Using simprofile, you can report the number of delay events being generated during simulation. A delay event is usually generated by Verilog expressions, which contains delay controls. For example:

```
#1;  
a = #1 0;  
a <= #1 1;  
#1 -> e;
```

However, after VCS optimizations, delay controls might be removed or merged. Therefore, the number of delay events can be less than the delay controls executed. Instrumentation on each delay event can be significant and therefore the performance overhead caused by profiling overhead is large.

Note:

Zero-delay (#0) event is not supported.

To enable, perform the following:

1. Specify the following compile option:

```
-simprofile
```

2. Specify the following runtime option:

```
-simprofile delay
```

To view the profiling report in a web browser, such as Firefox, specify the following:

```
firefox profileReport.html
```

To view the profiling report, which enables you to view the source line information, specify the following:

```
profrpt -firefox profileReport.html
```

The profiling report displays information in views. To view the source code, choose the **Construct** view. Choose the construct. For example, in this figure, the highlighted link is clicked.

Figure 2 Profiling Report: Counter Time Construct View

Counter Time Construct View			
Name	Count	Percentage	
Construct [link]	3743	80.01 %	
Construct [link]	2463	52.68 %	
Construct [link]	1234	26.20 %	
Construct [link]	2463	52.68 %	
Construct [link]	2463	52.68 %	
total	4671	100.00 %	

To view the source code, scroll down and click the **Source Information** link.

Figure 3 Profiling Report: Construct Information

Construct Information	
Construct Name	HzName (1e 7)
Time	2468
Construct Type	Initial
Parent Node	...
Source Information	View Source Code

The source code is displayed in a pop-up window. The counter of delay events are annotated to each source code line, as shown in the following screenshot.

Figure 4 Profiling Report: Source Code Information

```
01 module top;
02 initial begin
1234 20.00% 03     repeat (1234) #1;
04     $display("block end");
05 end
06 reg a;
07 initial begin
1234 20.00% 08     repeat (1234) a = #1 $time%2;
09     $display("a=", a);
1234 20.00% 10     repeat (1234) a = #1 $time%2;
11     $display("a=", a);
12 end
13 reg b;
2468 39.99% 14 always @(a) begin
15     b <= #1 a;
16 end
1 0.02% 17 initial begin
18     #3000;
19     $display("b=", b);
20 end
21 endmodule
22
```

8

Timing Analysis

Simulation Acceleration Transactor Code developers allow themselves a certain level of flexibility in code being mapped to ZeBu HW. Alternatively, RTL code developers for code mapped to silicon, have strict coding guidelines to meet size, power, timing and P&R constraints. However, Simulation Acceleration Transactor code mapped to HW still needs to meet timing constraints in order to function correctly. Therefore, understanding timing reports is important. Typically, performance of HW is determined by the `zTime`, which is the maximum frequency achieved by the driver clock.

To understand driver clock timing, critical paths timing reports can be analyzed. Then, coding changes to achieve a better `zTime`, which may lead to better performance, can be made.

For more information, see the following subsections:

- [UTF File Changes](#)
- [Viewing the zTime Report for Critical Output Routing Data Paths](#)
- [Analyzing a Critical Path Using zTime](#)

UTF File Changes

The usage of complex math operators (for example, modulo, divide, multiply, and so on) in the Testbench code mapped to HW that gets mapped by default to DSP Xilinx FPGA components, are not timed correctly. Therefore, it is recommended to disable DSP mapping by providing the following command in the UTF file:

```
optimization -dsp_mult_threshold 1024 -dsp_limit 1
```

This makes sure that operators using operands with less than 1024 bits are not mapped to a DSP Xilinx component.

It is recommended that you use the more accurate timing flow, which is called the SDF flow. This flow uses SDF timing arcs from the Xilinx component library to achieve a more accurate flow.

Add the following compile options to the UTF file:

```
# Enables SDF timing
timing_analysis -post_fpga BACK_ANNOTATED
# Enable better timing engines around memory timing analysis
timing_analysis -advanced_command {improved_memory_timing on}
ztopbuild -advanced_command {enable zmem_clock_domain_instrument}
```

The first command enables SDF timing. SDF timing provides additional timing information over traditional timing flows that work without using this command. With SDF timing, typically the `zTime` is lower than the `zTime` without SDF timing. With this, the timing is more conservative and runtime performance is worse. This does not mean your design will not run functionally correct if you do not apply the `sdf` UTF command.

For the timing analysis flow described in this section, SDF is required. After you have completed the timing analysis and timing fixes, you should remove this command to see if you can get a functional model without using SDF.

The second and third command are recommended because they enable better timing engines around memory timing analysis.

Viewing the zTime Report for Critical Output Routing Data Paths

After the `zCui` compile is complete, you can read and analyze a `zTime` report at the system level. Since the `zTime` report is in HTML format, open an HTML browser, such as Mozilla FireFox, in the `<zcu.work>/backend_default` directory.

To open a file, select a file URL with the full path in your Linux environment, using the `file://` directive. For example:

```
file://<full-path-to-your-html-file>
```

Perform the following steps to open the `zTime` report and view the critical output routing data paths:

1. Open the `zTime_fpga.html` report in an HTML browser, such as Mozilla FireFox. For example:

```
file:///home/projects/zcu.work/backend_default/zTime_fpga.html
```

The **zTime Report: Index** page appears.

zTime Report: Index

Date: Tue Jan 15 23:17:32 2019

[Filter paths](#)

[Routing data paths](#)

- Click **Routing data paths**. The **zTime Report: Critical Output Routing Data Paths** report appears. This report lists the critical paths, starting from the longest critical path.

Figure 5 zTime Report: Critical Output Routing Data Paths

zTime Report: critical output routing data paths
Date: Tue Jan 15 23:17:31 2019

Slack	Required Time	Delay	Fpga	From	To
0 ns	1566 ns (1 driver clock)	1566 ns	3 (2)	U0_M0_F11.zcbsplit_top_2345	U0_M0_F2.zcbsplit_top_638
0 ns	1566 ns (1 driver clock)	1566 ns	3 (2)	U0_M0_F11.zcbsplit_top_2345	U0_M0_F2.zcbsplit_top_638
176 ns	1566 ns (1 driver clock)	1389 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[6401]	U0_M0_F1.zcbsplit_p_zpor
176 ns	1566 ns (1 driver clock)	1389 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[6401]	U0_M0_F1.zcbsplit_p_zpor
189 ns	1566 ns (1 driver clock)	1377 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[1970]	U0_M0_F0.zcbsplit_p_zpor
189 ns	1566 ns (1 driver clock)	1377 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[1970]	U0_M0_F0.zcbsplit_p_zpor
216 ns	1566 ns (1 driver clock)	1350 ns	3 (2)	U0_M0_F11.zcbsplit_top_1701	U0_M0_F3.zcbsplit_top_850
216 ns	1566 ns (1 driver clock)	1350 ns	3 (2)	U0_M0_F11.zcbsplit_top_1701	U0_M0_F3.zcbsplit_top_850
880 ns	1566 ns (1 driver clock)	686 ns	3 (2)	U0_M0_F11.zcbsplit_top_3800	U0_M0_F4.zcbsplit_top_271
996 ns	1566 ns (1 driver clock)	570 ns	11 (7)	U0_M3_F10.ztbsplit_p_zport_w	U0_M0_F2.ztbsplit_top_117

The following table provides a description of the columns in the *zTime Report: Critical Output Routing Data Paths* report.

Table 7 zTime Report: Critical Output Routing Data Paths Column Descriptions

Column Name	Description
Slack	The amount of time in nanoseconds from the worst critical path. In the screenshot above, note that there is a sudden jump from 216 ns to 880 ns. This means that you get significant gains if you resolve the first 7 critical paths.
Required Time	This specifies how much delay is required for the most critical path. Our goal is to reduce this delay as much as possible. This column also specifies how many driver clock periods are used to complete the path. Note that the delay is based on the worst path for all the other paths.
Delay	This specifies the delay for the path.

Table 7 *zTime Report: Critical Output Routing Data Paths Column Descriptions*
(Continued)

Column Name	Description
Fpga	This specifies the number of hops required for the path. The number of hops is the number of FPGAs the path exits an FPGA and enters a different FPGA. In the parenthesis is the overall number of different FPGAs this path goes through. In this example, there is a path with 11 hops; this should be analyzed and reduced.
From	This is the source of the path. Refer to other documentation on how to use zBrowser to find this signal in the netlist, then how to find the relevant RTL signals in the Model's code.
To	This is the target of the path.

Analyzing a Critical Path Using zTime

After opening the **zTime Report: Critical Output Routing Data Paths** report, you can view the critical paths, starting from the longest critical path.

To analyze a path, in the **Slack** column, click the hyperlink for the path you want to analyze. In this section, the fifth path is analyzed, as highlighted in the following screenshot.

Figure 6 *zTime Report: Analyzing Critical Output Routing Data Paths*

zTime Report: critical output routing data paths
Date: Tue Jan 15 23:17:31 2019

Slack	Required Time	Delay	Fpga	From	To
0 ns	1566 ns (1 driver clock)	1566 ns	3 (2)	U0_M0_F11.zcbsplit_top_2345	U0_M0_F2.zcbsplit_top_638
0 ns	1566 ns (1 driver clock)	1566 ns	3 (2)	U0_M0_F11.zcbsplit_top_2345	U0_M0_F2.zcbsplit_top_638
176 ns	1566 ns (1 driver clock)	1389 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[6401]	U0_M0_F1.zcbsplit_p_zpor
176 ns	1566 ns (1 driver clock)	1389 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[6401]	U0_M0_F1.zcbsplit_p_zpor
189 ns	1566 ns (1 driver clock)	1377 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[1970]	U0_M0_F0.zcbsplit_p_zpor
189 ns	1566 ns (1 driver clock)	1377 ns	3 (2)	U0_M0_F11.zcbsplit_p_in[1970]	U0_M0_F0.zcbsplit_p_zpor
216 ns	1566 ns (1 driver clock)	1350 ns	3 (2)	U0_M0_F11.zcbsplit_top_1701	U0_M0_F3.zcbsplit_top_850
216 ns	1566 ns (1 driver clock)	1350 ns	3 (2)	U0_M0_F11.zcbsplit_top_1701	U0_M0_F3.zcbsplit_top_850
880 ns	1566 ns (1 driver clock)	686 ns	3 (2)	U0_M0_F11.zcbsplit_top_3800	U0_M0_F4.zcbsplit_top_271
996 ns	1566 ns (1 driver clock)	570 ns	11 (7)	U0_M3_F10.ztbsplit_p_zpor_w	U0_M0_F2.ztbsplit_top_117

The following page appears:

Figure 7 zTime Report: Segments in Critical Path

Fpga	Delay	Arrival	XDR	XTYPE	zCore	Port
Internal					DriverClock domain	
U0/M0/F11	17 ns	17 ns			part_U0_M0	U0_M0_F11/zcbsplt_p_in[1970]
	148 ns		144	LVDS		part_U0_M0@UC
U0/M0/F05	8 ns	173 ns				
	135 ns		128	LVDS		
U0/M0/F00		308 ns			part_U0_M0	U0_M0_F0/zcbsplt_p_in[1970]
Internal	1069 ns	1377 ns			DutClock domain	

This page provides the different segments in the path. Each segment involves exiting an FPGA (starting FPGA) and entering a different FPGA (ending FPGA). The following table provides a description of the columns.

Table 8 zTime Report: Column Description of Segments in Critical Path

Column	Description
Fpga	Name of the FPGA from where the current segment begins
Delay	The time spent in a segment
Arrival	The absolute time the segment started
Port	Source signal of the segment

In the preceding screenshot, notice that the time in the FPGA specified as **internal** is unexpectedly large and sometimes seems that it repeats for all segments going from FPGA x to FPGA y. Timing in ZeBu is conservative. For each segment, the tool accounts for the most critical path known in the FPGAs in the respective segment.

The only way to improve on timing related to FPGAs is to produce a critical path report for the specific FPGA. Go to the following location to view FPGA timing reports at FPGA-compile level:

```
<path-to-zcui.work>/backend_default/U*/M*/F*/fpga_reports/vivado/timing_optimized_hold.rpt
```

Analysis of these critical paths is beyond the scope of this document.

9

Using Time Decoupled Mode for Performance Improvements

Simulation Acceleration enables you to mark functions, which only have inputs, so that these functions are not monitored or controlled by Simulation Acceleration transactors. This is known as time decoupling of functions, which improves performance. This technique does not affect the timestamp synchronization.

For more information, see the following subsections:

- [Consideration for Using Time Decoupling](#)
- [Use Model for the Time Decouple Mode](#)
- [Limitations of Time Decouple Mode](#)

Consideration for Using Time Decoupling

Before using time decoupling, review the following considerations:

- With the time decoupled mode, the result from run-to-run might be different. When the waveform outputting (FWC/QiWC/read-back) is enabled, the discrepancy is significant between runs that use time coupled mode and those that do not.
- The timestamp is non-deterministic in the following cases:
 - If the criterion to finish the emulation does not rely on a `#delay`, the timestamps between the testbench and HW at the end of emulation might be different and non-deterministic.
 - If a decoupled function on the testbench side called by HW triggers a primary inputs (PI), the timestamp at which this PI arrives from the HW side is non-deterministic.

From a HW perspective, PI refers to the data transmitted from the testbench to DUT by anyone of following ways: Port Connections, XMR, and Functions Calls.

Port Connection	XMR	Function Call
<pre>Module Testbench; Wire tb_sig; DUT u_DUT(.sig_in (tb_sig)); Endmodule</pre>	<pre>Assign Testbench.u_DUT.dut_sig = tb_sig;</pre>	<pre>Testbench.u_DUT.fun c();</pre>

- If a HW function takes multiple driver clocks to complete execution in the same delta time, it is not recommended to be marked as time decoupled. For the following snippet, ZEMI3 will stop `Cclock` until the function completes execution. Therefore, the performance cannot be improved in the time decoupled mode.

```
function void initialize_memory (input bit [31:0] init_val);
  int i = 0;
  while (i < 8192) begin
    mem[i] = init_val; // mem is a 2-D memory
    i++;
  end
endfunction
```

Use Model for the Time Decouple Mode

To use the time decouple mode, see the following subsections:

- [Compile-Time Use Model](#)
- [Runtime Use Model](#)

Compile-Time Use Model

The following methods are available to compile functions in the time coupled mode.

- [Automatically Infer All Input-Only Functions](#)
- [Manually Specify Input-Only Functions](#)

Note:

Only functions, which are compiled in the time coupled mode, are configurable at runtime time decoupled mode.

Automatically Infer All Input-Only Functions

When you use this method, Simulation Acceleration compiles *all* functions, which have only inputs ports, in the time decoupled mode. These functions are then configurable during runtime.

To automatically infer and compile all functions, which have only input ports, specify the following switch:

```
-simxl=time_decouple+input_only_function
```

Manually Specify Input-Only Functions

If you want to select specific input-only functions to compile in the time coupled mode, you need to first make RTL changes to the functions. In the RTL code, use the `simxl_time_decouple` attribute on a function that has input ports only. You can then control these functions during runtime in time coupled mode.

In the following RTL code snippet, the `user_func` function is compiled in time decoupled mode and configurable at runtime.

```
(* simxl_time_decouple=1 *)  
function void user_func(input arg);  
endfunction
```

To make sure that a function, which contains only input ports, is neither compiled in time decoupled mode, and therefore not configurable at runtime, set the `simxl_time_decouple` attribute to 0.

To compile the functions in the time decoupled mode, pass the following compile-time switch during compilation using `simv`:

```
-simxl=time_decouple
```

Runtime Use Model

The `-simxl=time_decouple` runtime command enables the time decoupled mode on all the functions that are compiled for time decoupled mode. To enable the time decouple mode work normally, set the `SNPS_SIMXL_DISABLE_CCC` environment variable to 1.

To facilitate debugging during runtime, you can overwrite the default settings specified at compile time. To do this, pass a configuration file by specifying the following:

```
-simxl=time_decouple_cfg:<file path>
```

Note:

If `-simxl=time_decouple+time_decouple_cfg:<>` is passed, the configuration file overrides the functions enabled by the `-simxl=time_decouple` command.

The runtime configuration file contains entries for each function call that you want to run in the time decoupled mode. The format of each entry is as follows:

```
<module_name::tf_name>, <time_decouple_val 0|1>, <list_of_ids  
corresponding to the tf>
```

Simulation Acceleration outputs a sample runtime configuration file, which you can copy, edit, and use in your runtime. The location of this file is as follows:

```
zcui.work/default_backend/time_decouple_cfg.csv
```

In addition, see the following file for detailed information on each function call. Information, such as function call ID with reference to the compilation and information about function call scope and location, is available. The location of this file is as follows:

```
zcui.work/default_backend/time_decouple_tf_call_info.json
```

Limitations of Time Decouple Mode

The following limitations exist with this Simulation Acceleration feature:

- Functions with output ports are not supported for the time decouple mode. VCS reports an NYI error at compile time.
- Tasks are not supported for the time decoupled mode.
- Under the compile time switch, `-simx1=time_decouple`, the `$display` family of system tasks and UVM calls do not have time decouple support.

10

Enabling zDPI in Simulation Acceleration

ZeBu supports a subset of DPI-C standard through the zDPI, where only the data is going from HW to SW. Therefore, if the HW design contains import DPI functions with only inputs as arguments and no return value, DPI functions can process using zDPI.

You can enable zDPI support in Simulation Acceleration in both V2Z Mode and V2VX Mode.

For more information, see the following topics:

- [Use Model for Enabling zDPI in the V2Z Flow](#)
- [Use Model for Enabling zDPI in the V2VX Flow](#)
- [Support to Call Export DPI Function Calls](#)
- [Limitations of zDPI Support in Simulation Acceleration](#)

Use Model for Enabling zDPI in the V2Z Flow

To enable zDPI in the V2Z Mode, you need to make compilation and runtime changes.

Refer to the following topics

- [Compile-Time Use Model](#)
- [Runtime Use Model](#)

Compile-Time Use Model

To enable zDPI, specify the following UTF command:

```
simx1_enable_zdpi -all <0|1> | -modules {modules list}
```

Where:

- `-all`: Use to enable zDPI for all modules.
- `-modules`: Use to enable zDPI for specific modules. This option takes a space-separated list of modules. For the specified modules, Simulation Acceleration skips its own DPI processing and the specified modules get processed as zDPI by Unified FrontEnd (UFE).

In addition, specify the following UTF command:

```
dpi_synthesis -enable ALL
```

Runtime Use Model

By default, zDPI is disabled at runtime. To enable zDPI at runtime, specify the following option:

```
-simxl=zdpi:path_to_so
```

If the zDPI allows multiple shared objects, specify the `-simxl=zdpi:path_to_so` option for each object.

Use Model for Enabling zDPI in the V2VX Flow

In the V2VX Mode, use the following option to enable zDPI globally:

```
-simxl=enable_zdpi
```

To enable zDPI for specific modules, use the plus sign (+) to specify the module names. For example, the following command specification enables zDPI in module1 and module2:

```
-simxl=enable_zdpi+module1+module2
```

Support to Call Export DPI Function Calls

An import DPI function calling an export DPI function in a DUT module during simulation is now supported by adding the following switch to VCS compilation command line:

```
-Xhwcosim=import_dpi_func_to_task
```

SimXL rewrites all import context DPI functions when the switch is added.

Limitations of zDPI Support in Simulation Acceleration

The following limitations are applicable to all imported DPI in zDPI modules:

- Only input ports are supported. The following are not supported: output ports, inout ports, and ref ports
- The import DPI should not have a return value
- Import DPI tasks are not allowed because they may potentially be blocking

If these guidelines are not followed and zDPI is enabled on modules, an error is reported.

However, if zDPI enabled globally, instead of module-wise, a message is not reported and zDPI might not be able to work correctly.

11

Estimating DUT Power Requirements Using Power Profiling

The Power Profiling feature in Simulation Acceleration enables you to estimate the power requirements of the DUT, which resides in the hardware. During Simulation Acceleration runtime, power profiling is performed using the `wap` UCLI command.

Enabling Weighted Activity Profile (WAP) Power Profiling

To enable WAP power profiling, specify the following UTF command at compile time:

```
power_profile -weight_profile <file>
```

Where, `-weight_profile <file>` specifies the power weight file generated by SpyGlass.

Use Model of the `wap` Command

In Simulation Acceleration, you specify `wap` UCLI commands in runtime. To perform power profiling in Simulation Acceleration, you can specify a UCLI script or enter commands manually in UCLI mode.

The `wap` command is used with the following options:

```
wap
-defined
-init
-close
-list atc_threshold|enabled|threshold|weight -domain <name>
-output_dir [-set <path>]
-total weight|threshold|atc_threshold
-mode [-per_bucket|-total_bucket]
-bucket -config_file [-default|-import <path>|-export <path>]
-bucket -name <name> -enable|-disable
-bucket -name <name> -threshold|-atc_threshold|-weight <value>
-bucket -name <name> -create_from <bucket_list>
-bucket -name <name> -delete
-dump -enable|-disable
```

Where:

- `-defined`: Used to verify whether power profiling is available in the session.
- `-init`: Used to initialize power profiling.
- `-close`: Used to close and flush all power profiling operations.
- `-list atc_threshold|enabled|threshold|weight -domain <name>`: Used to list property values of domains.
- `-output_dir [-set <path>]`: Used to retrieve the current output path of power profiling or set the output path of power profiling.
- `-total weight|threshold|atc_threshold`: Used to retrieve the total value of a specific property.
- `-mode [-per_bucket|-total_bucket]`: Used to retrieve the current WAP run mode. You can set the run mode by specifying either `-per_bucket` or `-total_bucket`.
- `-bucket -config_file [-default|-import <path>|-export <path>]`: Used to set the wap configuration values to default values. Also, used to set the configuration values from an imported configuration file. In addition, this option is used to export current configuration values to a file.
- `-bucket -name <name> -enable|-disable`: Used to enable or disable a bucket for power profiling.
- `-bucket -name <name> -threshold|-atc_threshold|-weight <value>`: Used to set the property values of a given bucket to a given value.
- `-bucket -name <name> -create_from <bucket_list>`: Used to create a new bucket from a list of existing buckets.
- `-bucket -name <name> -delete`: Used to delete a bucket from power profiler.
- `-dump -enable|-disable`: Used to enable and disable power profiler dumping operations.

Sample UCLI Script Using the WAP Command Specifications

The following is a sample UCLI script that has wap command specifications for power profiling.

```
==== ucli.tcl=====
wap -defined
wap -init
wap -dump -enable
run 1000
wap -dump -disable
```

```
wap -close  
exit
```